
mahotas Documentation

Release 1.4.3+git

Luis Pedro Coelho

Jun 14, 2017

Contents

1 Examples	3
2 Full Documentation Contents	5
2.1 How To Install Mahotas	5
2.2 Finding Wally	7
2.3 Labeled Image Functions	9
2.4 Thresholding	22
2.5 Wavelet Transforms	28
2.6 Distance Transform	35
2.7 Polygon Utilities	40
2.8 Features	42
2.9 Local Binary Patterns	43
2.10 Speeded-Up Robust Features	44
2.11 Implementing SURF-ref With Mahotas	48
2.12 Morphological Operators	50
2.13 Color Space Conversions	55
2.14 Input/Output with Mahotas	57
2.15 Tutorial: Classification Using Mahotas	58
2.16 Tutorial: Extended Depth of Field	59
2.17 mahotas-features.py	61
2.18 Frequently Asked Questions	62
2.19 Mahotas Internals	64
2.20 The Why of mahotas	65
2.21 Contributing	67
2.22 Possible Tasks	67
2.23 History	68
2.24 Full API Documentation	77
3 Indices and tables	105
Bibliography	107
Python Module Index	109

Note: If you are using mahotas in a scientific publication, please cite:

Coelho, L.P. 2013. Mahotas: Open source software for scriptable computer vision. Journal of Open Research Software 1(1):e3, DOI: <http://dx.doi.org/10.5334/jors.ac>

Mahotas is a computer vision and image processing library for Python.

It includes many algorithms implemented in C++ for speed while operating in numpy arrays and with a very clean Python interface.

Mahotas currently has over 100 functions for image processing and computer vision and it keeps growing. Some examples of mahotas functionality:

- [watershed](#)
- convex points calculations.
- [hit & miss. thinning](#)
- Zernike & Haralick, [local binary patterns](#), and TAS features.
- morphological processing
- Speeded-Up Robust Features (SURF), a form of local features
- [thresholding](#)
- convolution.
- Sobel edge detection.

The release schedule is roughly one release every few months and each release brings new functionality and improved performance. The interface is very stable, though, and code written using a version of mahotas from years back will work just fine in the current version, except it will be faster (some interfaces are deprecated and will be removed after a few years, but in the meanwhile, you only get a warning).

Bug reports with test cases typically get fixed in 24 hours.

See also:

[mahotas-imread](#) is side project which includes code to read/write images to files

CHAPTER 1

Examples

This is a simple example of loading a file (called *test.jpeg*) and calling *watershed* using above threshold regions as a seed (we use Otsu to define threshold).

```
import numpy as np
import mahotas
import pylab

img = mahotas.imread('test.jpeg')
T_otsu = mahotas.thresholding.otsu(img)
seeds,_ = mahotas.label(img > T_otsu)
labeled = mahotas.cwatershed(img.max() - img, seeds)


pylab.imshow(labeled)
```

Computing a distance transform is easy too:

```
import pylab as p
import numpy as np
import mahotas

f = np.ones((256,256), bool)
f[200:,240:] = False
f[128:144,32:48] = False
# f is basically True with the exception of two islands: one in the lower-right
# corner, another, middle-left

dmap = mahotas.distance(f)
p.imshow(dmap)
p.show()
```



CHAPTER 2

Full Documentation Contents

[Jump to detailed API Documentation](#)

How To Install Mahotas

The simplest way to install mahotas is to use `conda`.

If you have `conda` installed, you can install mahotas using the following pair of commands:

```
conda config --add channels conda-forge  
conda install mahotas
```

This relies on the `conda-forge` project, which builds packages for all major environments (Linux, Mac OS X, and Windows). If you do not want to permanently add the `conda-forge` channel to your `conda` configuration, you can also install just mahotas with:

```
conda install -c https://conda.anaconda.org/conda-forge mahotas
```

From source

You can get the released version using `pip`:

```
pip install mahotas
```

If you prefer, you can download the source from [PyPI](#) and run:

```
python setup.py install
```

You will need to have `numpy` and a C++ compiler.

Visual Studio

For compiling from source in Visual Studio, use:

```
python setup.py build_ext -c msvc  
python setup.py install
```

Bleeding Edge (Development)

Development happens on [github](#). You can get the development source there. Watch out that *these versions are more likely to have problems*.

Packaged Versions

On Windows

On Windows, Christoph Gohlke does an excellent job maintaining binary packages of mahotas (and several other packages).

WinPython

WinPython ships with mahotas as a standard package

Python(x, y)

If you use Python(x, y), which is often a good solution, then you probably have it already as mahotas is a standard plugin.

Enthought Canopy

Since May 2015, Enthought's Canopy Package Index includes mahotas.

FreeBSD

Mahotas is available for FreeBSD as [graphics/mahotas](#).

MacPorts

For Macports, mahotas is available as [py27-mahotas](#).

conda

Mahotas is not a part of standard conda packages, but on 64 bit Linux, you can get it [from this repository](#) with:

```
conda install -c https://conda.binstar.org/luispedro mahotas
```

Frugalware Linux


Mahotas is available as `python-mahotas`.

Finding Wally

This was originally an [answer on stackoverflow](#) We can use it as a simple tutorial example.

The problem is to find Wally (who goes by Waldo in the US) in the following image:

```
from pylab import imshow, show
import mahotas
import mahotas.demos
wally = mahotas.demos.load('Wally')
imshow(wally)
show()
```



From October 11 2013 onwards (version 1.0.4 or later), you can get the Wally image from mahotas as:

```
import mahotas.demos
wally = mahotas.demos.load('Wally')
```

Can you see him?

```
wffloat = wally.astype(float)
r,g,b = wffloat.transpose((2,0,1))
```

Split into red, green, and blue channels. It's better to use floating point arithmetic below, so we convert at the top.

```
w = wffloat.mean(2)
```

w is the white channel.

```
pattern = np.ones((24,16), float)
for i in xrange(2):
    pattern[i::4] = -1
```

Build up a pattern of +1,+1,-1,-1 on the vertical axis. This is Wally's shirt.

```
v = mahotas.convolve(r-w, pattern)
```

Convolve with red minus white. This will give a strong response where the shirt is.


```
mask = (v == v.max())
mask = mahotas.dilate(mask, np.ones((48,24)))
```

Look for the maximum value and dilate it to make it visible. Now, we tone down the whole image, except the region of interest:

```
wally -= .8*wally * ~mask[:, :, None]
```

And we get the following:

```
wffloat = wally.astype(float)
r,g,b = wffloat.transpose((2,0,1))
w = wffloat.mean(2)
pattern = np.ones((24,16), float)
for i in xrange(2):
    pattern[i::4] = -1
v = mahotas.convolve(r-w, pattern)
mask = (v == v.max())
mask = mahotas.dilate(mask, np.ones((48,24)))
np.subtract(wally, .8*wally * ~mask[:, :, None], out=wally, casting='unsafe')
imshow(wally)
show()
```



Labeled Image Functions

Labeled images are integer images where the values correspond to different regions. I.e., region 1 is all of the pixels which have value 1, region two is the pixels with value 2, and so on. By convention, **region 0 is the background and often handled differently**.

Labeling Images

New in version 0.6.5.


The first step is obtaining a labeled function from a binary function:

```
import mahotas as mh
import numpy as np
from pylab import imshow, show

regions = np.zeros((8,8), bool)

regions[:3,:3] = 1
regions[6:,6:] = 1
labeled, nr_objects = mh.label(regions)
```

```
imshow(labeled, interpolation='nearest')
show()
```



This results in an image with 3 values:

0. background, where the original image was 0
1. for the first region: (0:3, 0:3);
2. for the second region: (6:, 6:).

There is an extra argument to `label`: the structuring element, which defaults to a 3x3 cross (or, 4-neighbourhood). This defines what it means for two pixels to be in the same region. You can use 8-neighbourhoods by replacing it with a square:

```
labeled, nr_objects = mh.label(regions, np.ones((3,3), bool))
```

We can now collect a few statistics on the labeled regions. For example, how big are they?

```
sizes = mh.labeled.labeled_size(labeled)
print('Background size', sizes[0])
print('Size of first region: {}'.format(sizes[1]))
```

This size is measured simply as the number of pixels in each region. We can instead measure the total weight in each area:

```
array = np.random.random_sample(regions.shape)
sums = mh.labeled_sum(array, labeled)
print('Sum of first region: {}'.format(sums[1]))
```

Filtering Regions


New in version 0.9.6: `remove_regions` & `relabel` were added.

Here is a slightly more complex example. The full code is in the `demos` directory as `nuclear.py`. We are going to use this image, a fluorescent microscopy image from a [nuclear segmentation benchmark](#)

This image is available as `mahotas.demos.nuclear_image()`

```
import mahotas as mh
import mahotas.demos
import numpy as np
from pylab import imshow, show

f = mh.demos.nuclear_image()
f = f[:, :, 0]
imshow(f)
show()
```




First we perform a bit of Gaussian filtering and thresholding:

```
f = mh.gaussian_filter(f, 4)
f = (f > f.mean())
```

(Without the Gaussian filter, the resulting thresholded image has very noisy edges. You can get the image in the demos/ directory and try it out.)


```
f = mh.gaussian_filter(f, 4)
f = (f > f.mean())
imshow(f)
show()
```



Labeling gets us all of the nuclei:

```
labeled, n_nucleus = mh.label(f)
print('Found {} nuclei.'.format(n_nucleus))
```


```
labeled, n_nucleus = mh.label(f)
print('Found {} nuclei.'.format(n_nucleus))
imshow(labeled)
show()
```



42 nuclei were found. None were missed, but, unfortunately, we also get some aggregates. In this case, we are going to assume that we wanted to perform some measurements on the real nuclei, but are willing to filter out anything that is not a complete nucleus or that is a lump on nuclei. So we measure sizes and filter:

```
sizes = mh.labeled.labeled_size(labeled)
too_big = np.where(sizes > 10000)
labeled = mh.labeled.remove_regions(labeled, too_big)
```


```
sizes = mh.labeled.labeled_size(labeled)
too_big = np.where(sizes > 10000)
labeled = mh.labeled.remove_regions(labeled, too_big)
imshow(labeled)
show()
```



We can also remove the region touching the border:

```
labeled = mh.labeled.remove_bordering(labeled)
```

```
labeled = mh.labeled.remove_bordering(labeled)
imshow(labeled)
show()
```




This array, `labeled` now has values in the range 0 to `n_nucleus`, but with some values missing (e.g., if region 7 was one of the ones touching the border, then 7 is not used in the labeling). We can `relabel` to get a cleaner version:

```
relabelled, n_left = mh.labeled.relabel(labeled)
print('After filtering and relabeling, there are {} nuclei left.'.format(n_left))
```

Now, we have 24 nuclei and `relabelled` goes from 0 (background) to 24.

```
relabelled, n_left = mh.labeled.relabel(labeled)
print('After filtering and relabeling, there are {} nuclei left.'.format(n_left))
imshow(relabelled)
show()
```



In mahotas after version 1.4, we can even make many of the same operations with a single call to `mh.labeled.filter_labeled`:

```
relabelled, n_left = mh.labeled.filter_labeled(labeled, remove_bordering=True, max_size=10000)
```

Borders

A border pixel is one where there is more than one region in its neighbourhood (one of those regions can be the background).

You can retrieve border pixels with either the `borders()` function, which gets all the borders or the `border()` (note the singular) which gets only the border between a single pair of regions. As usual, what neighbour means is defined by a structuring element, defaulting to a 3x3 cross.

API Documentation

The `mahotas.labeled` submodule contains the functions mentioned above. `label()` is also available as `mahotas.label`.

`mahotas.labeled.bbox(f, as_slice=False)`
Bounding boxes of all objects in a labeled array.

After:

```
bboxes = mh.labeled.bbox(f)
```

`bboxes[34]` will contain the bounding box of (`f == 34`).

Parameters `f`: integer ndarray

as_slice : boolean, optional

Whether to return slice objects instead of integer coordinates (default: False).

Returns `bboxes` : ndarray

See also:

mh.bbox the binary version of this function

`mahotas.labeled.borders(labeled, Bc={3x3 cross}, out={np.zeros(labeled.shape, bool)})`

Compute border pixels

A pixel is on a border if it has value i and a pixel in its neighbourhood (defined by Bc) has value j , with $i \neq j$.

Parameters `labeled` : ndarray of integer type

input labeled array

Bc : structure element, optional

out : ndarray of same shape as `labeled`, dtype=bool, optional

where to store the output. If `None`, a new array is allocated

mode : {'reflect', 'nearest', 'wrap', 'mirror', 'constant' [default], 'ignore'}

How to handle borders

Returns `border_img` : boolean ndarray

Pixels are True exactly where there is a border in `labeled`

`mahotas.labeled.border(labeled, i, j, Bc={3x3 cross}, out={np.zeros(labeled.shape, bool)}, always_return=True)`

Compute the border region between i and j regions.

A pixel is on the border if it has value i (or j) and a pixel in its neighbourhood (defined by Bc) has value j (or i).

Parameters `labeled` : ndarray of integer type

input labeled array

i : integer

j : integer

Bc : structure element, optional

out : ndarray of same shape as `labeled`, dtype=bool, optional

where to store the output. If `None`, a new array is allocated

always_return : bool, optional

if false, then, in the case where there is no pixel on the border, returns `None`. Otherwise (the default), it always returns an array even if it is empty.

Returns `border_img` : boolean ndarray

Pixels are True exactly where there is a border between i and j in `labeled`

`mahotas.labeled.bwperim(bw, n=4)`

Find the perimeter of objects in binary images.

A pixel is part of an object perimeter if its value is one and there is at least one zero-valued pixel in its neighborhood.

By default the neighborhood of a pixel is 4 nearest pixels, but if n is set to 8 the 8 nearest pixels will be considered.

Parameters `bw` : ndarray

A black-and-white image (any other image will be converted to black & white)

n : int, optional

Connectivity. Must be 4 or 8 (default: 4)

mode : {'reflect', 'nearest', 'wrap', 'mirror', 'constant' [default], 'ignore'}

How to handle borders

Returns `perim` : ndarray

A boolean image

See also:

borders function This is a more generic function

`mahotas.labeled.filter_labeled(labeled, remove_bordering=False, min_size=None, max_size=None)`

Filter labeled regions based on a series of conditions

New in version 1.4.1.

Parameters `labeled` : labeled array

`remove_bordering` : bool, optional

whether to remove regions that touch the border

`min_size` : int, optional

Minimum size (in pixels) of objects to keep (default is no minimum)

`max_size` : int, optional

Maximum size (in pixels) of objects to keep (default is no maximum)

Returns `filtered` : labeled array

`nr` : int

number of new labels

`mahotas.labeled.label(array, Bc={3x3 cross}, output={new array})`

Label the array, which is interpreted as a binary array

This is also called *connected component labeled*, where the connectivity is defined by the structuring element `Bc`.

See: http://en.wikipedia.org/wiki/Connected-component_labeling

Parameters `array` : ndarray

This will be interpreted as binary array

`Bc` : ndarray, optional

This is the structuring element to use

out : ndarray, optional

Output array. Must be a C-array, of type np.int32

Returns labeled : ndarray

Labeled result

nr_objects : int

Number of objects

`mahotas.labeled.labeled_sum(array, labeled, minlength=None)`

Labeled sum. sum will be an array of size labeled.max() + 1, where sum[i] is equal to np.sum(array[labeled == i]).

Parameters array : ndarray of any type

labeled : int ndarray

Label map. This is the same type as returned from `mahotas.label()`

minlength : int, optional

Minimum size of return array. If labeled has fewer than minlength regions, 0s are added to the result. (optional)

Returns sums : 1-d ndarray of array.dtype

`mahotas.labeled.labeled_max(array, labeled, minlength=None)`

Labeled minimum. mins will be an array of size labeled.max() + 1, where mins[i] is equal to np.min(array[labeled == i]).

Parameters array : ndarray of any type

labeled : int ndarray

Label map. This is the same type as returned from `mahotas.label()`

Returns mins : 1-d ndarray of array.dtype

`mahotas.labeled.labeled_size(labeled)`

Equivalent to:

```
for i in range(...):
    sizes[i] = np.sum(labeled == i)
```

but, naturally, much faster.

Parameters labeled : int ndarray

Returns sizes : 1-d ndarray of int

See also:

`mahotas.fullhistogram` almost same function by another name (the only

difference

`mahotas.labeled.relabel(labeled, inplace=False)`

Relabeling ensures that relabeled is a labeled image such that every label from 1 to relabeled.max() is used (0 is reserved for the background and is passed through).

Example:

```
labeled,n = label(some_binary_map)
for region in xrange(n):
    if not good_region(labeled, region + 1):
        # This deletes the region:
        labeled[labeled == (region + 1)] = 0
relabel(labeled, inplace=True)
```

Parameters `relabeled` : ndarray of int

A labeled array

inplace : boolean, optional

Whether to perform relabeling inplace, erasing the values in `labeled` (default: False)

Returns `relabeled`: ndarray

nr_objs : int

Number of objects

See also:

`label` function

`mahotas.labeled.is_same_labeling(labeled0, labeled1)`

Checks whether `labeled0` and `labeled1` represent the same labeling (i.e., whether they are the same except for a possible change of label values).

Note that the background (value 0) is treated differently. Namely

`is_same_labeling(a, b)` implies `np.all((a == 0) == (b == 0))`

Parameters `labeled0` : ndarray of int

A labeled array

`labeled1` : ndarray of int

A labeled array

Returns `same` : bool

True if the labelings passed as argument are equivalent

See also:

`label` function

`relabel` function

`mahotas.labeled.perimeter(bwimage, n=4, mode="constant")`

Calculate total perimeter of all objects in binary image.

Parameters `bwimage` : array

binary image

`n` : int, optional

passed to `bwperim` as is

`mode` : str, optional

passed to `bwperim` as is

Returns `p` : float

total perimeter of all objects in binary image

See also:

`bwperim` function Finds the perimeter region

References

[R1]

`mahotas.labeled.remove_bordering(labeled, rsize=1, out={np.empty_like(im)})`

Remove objects that are touching the border.

Pass `labeled` as `out` to achieve in-place operation.

Parameters `labeled` : ndarray

Labeled array

`rsize` : int or tuple, optional

Minimum distance to the border (in Manhattan distance) to allow an object to survive.

May be int or tuple with len == `labeled.ndim`.

`out` : ndarray, optional

If `im` is passed as `out`, then it operates inline.

Returns `slabeled` : ndarray

Subset of `labeled`

`mahotas.labeled.remove_regions(labeled, regions, inplace=False)`

`removed = remove_regions(labeled, regions, inplace=False):`

Removes the regions in `regions`. If an elementwise `in` operator existed, this would be equivalent to the following:

```
labeled[ labeled element-wise-in regions ] = 0
```

This function **does not** relabel its arguments. You can use the `relabel` function for that:

```
removed = relabel(remove_regions(labeled, regions))
```

Or, saving one image allocation:

```
removed = relabel(remove_regions(labeled, regions), inplace=True)
```

This is the same, but reuses the memory in the relabeling operation.

Parameters `relabeled` : ndarray of int

A labeled array

`regions` : sequence of int

These regions will be removed

`inplace` : boolean, optional

Whether to perform removal inplace, erasing the values in `labeled` (default: False)

Returns removed : ndarray

See also:

relabel function After removing unnecessary regions, it is often a good idea to relabel your label image.

`mahotas.labeled.remove_regions_where(labeled, conditions, inplace=False)`

Remove regions based on a boolean array

A region is removed if `conditions[region-id]` evaluates true.

This function **does not** relabel its arguments. You can use the `relabel` function for that:

```
removed = relabel(remove_regions_where(labeled, conditions))
```

Or, saving one image allocation:

```
removed = relabel(remove_regions(labeled, conditions), inplace=True)
```

This is the same, but reuses the memory in the relabeling operation.

See also:

remove_regions function Variation of this function which uses integer indexing

Thresholding

The example in this section is present in the source under `mahotas/demos/thresholding.py`.

We start with an image, a grey-scale image:

```
import mahotas.demos
photo = mahotas.demos.load('luispedro')
photo = photo.astype(np.uint8)
```

Before Oct 2013, the `mahotas.demos.load` function did not exist and you needed to specify the path explicitly:


```
luispedro_image = '../../../../../mahotas/demos/data/luispedro.jpg'
photo = mahotas.imread(luispedro_image, as_grey=True)
```

The reason we convert to `np.uint8` is because `as_grey` returns floating point images (there are good reasons for this and good reasons against it, since it's easier to truncate than to go back, it returns `np.uint8`).

```
import mahotas
import mahotas.demos
import numpy as np
from pylab import imshow, gray, show
from os import path

photo = mahotas.demos.load('luispedro', as_grey=True)
photo = photo.astype(np.uint8)

gray()
imshow(photo)
show()
```



Thresholding functions have a trivial interface: they take an image and return a value. One of the most well-known thresholding methods is Otsu's method:


```
T_otsu = mahotas.otsu(photo)
print(T_otsu)
imshow(photo > T_otsu)
show()
```

prints 115.

```
import mahotas
import mahotas.demos
import numpy as np
from pylab import imshow, gray, show
from os import path

photo = mahotas.demos.load('luispedro', as_grey=True)
photo = photo.astype(np.uint8)

T_otsu = mahotas.otsu(photo)
print(T_otsu)
gray()
imshow(photo > T_otsu)
show()
```



An alternative is the Riddler-Calvard method:


```
T_rc = mahotas.rc(photo)
print(T_rc)
imshow(photo > T_rc)
show()
```

In this image, it prints almost the same as Otsu: 115.68. The thresholded image is exactly the same:

```
import mahotas
import mahotas.demos
import numpy as np
from pylab import imshow, gray, show
from os import path

photo = mahotas.demos.load('luispedro', as_grey=True)
photo = photo.astype(np.uint8)

T_rc = mahotas.rc(photo)
print(T_rc)
gray()
imshow(photo > T_rc)
show()
```



See also the labeled documentation which can be very helpful in combination with thresholding.

API Documentation

The `mahotas.thresholding` module contains the thresholding functions, but they are also available in the main `mahotas` namespace.

Thresholding Module

Thresholding functions.

These functions return **the numeric threshold**. In order to obtain a *thresholded image*, you can do the following:

```
T_otsu = mh.otsu(image)
binarized_image = (image > T_otsu)
```

Functions which have an `ignore_zeros` parameters will only consider non-zero pixels when computing the thresholding.

otsu() Otsu method

rc() Riddler-Calvard's method

bernsen Bernsen thresholding

gbernsen Generalized Bernsen thresholding

`mahotas.thresholding.otsu(img, ignore_zeros=False)`

Calculate a threshold according to the Otsu method.

Example:

```
import mahotas as mh
import mahotas.demos

im = mahotas.demos.nuclear_image()
# im is stored as RGB, let's convert to single 2D format:
im = im.max(2)

#Now, we compute Otsu:
t = mh.otsu(im)

# finally, we use the value to form a binary image:
bin = (im > t)
```

See Wikipedia for details on methods: http://en.wikipedia.org/wiki/Otsu's_method

Parameters `img` : an image as a numpy array.

This should be of an unsigned integer type.

`ignore_zeros` : Boolean

whether to ignore zero-valued pixels (default: False)

Returns `T` : integer

the threshold

`mahotas.thresholding.rc(img, ignore_zeros=False)`

Calculate a threshold according to the Riddler-Calvard method.

Example:

```
import mahotas as mh
import mahotas.demos

im = mahotas.demos.nuclear_image()
# im is stored as RGB, let's convert to single 2D format:
im = im.max(2)

#Now, we compute a threshold:
t = mh.rc(im)

# finally, we use the value to form a binary image:
bin = (im > t)
```

Parameters `img` : ndarray

Image of any type

`ignore_zeros` : boolean, optional


Whether to ignore zero valued pixels (default: False)

Returns `T` : float

threshold

`mahotas.thresholding.soft_threshold(f, tval)`

Soft threshold function:



Parameters `f` : ndarray

`tval` : scalar

Returns `thresholded` : ndarray

`mahotas.thresholding.bernsen(f, radius, contrast_threshold, gthresh={128})`

Bernsen local thresholding

Parameters `f` : ndarray

 input image

`radius` : integer

 radius of circle (to consider “local”)

`contrast_threshold` : integer

 contrast threshold

`gthresh` : numeric, optional

 global threshold to fall back in low contrast regions

Returns `thresholded` : binary ndarray

See also:

`gbernsen` function Generalised Bernsen thresholding

`mahotas.thresholding.gbernsen(f, se, contrast_threshold, gthresh)`

Generalised Bernsen local thresholding

Parameters `f` : ndarray

 input image

`se` : boolean ndarray

 structuring element to use for “locality”

`contrast_threshold` : integer

 contrast threshold

`gthresh` : numeric, optional

 global threshold to fall back in low contrast regions

Returns thresholded : binary ndarray

See also:

bernsen function Bernsen thresholding with a circular region

Wavelet Transforms

New in version 0.9.1: Wavelet functions were only added in version 0.9.1


We are going to use wavelets to transform an image so that most of its values are 0 (and otherwise small), but most of the signal is preserved.

The code for this tutorial is available from the source distribution as `mahotas/demos/wavelet_compression.py`.

We start by importing and loading our input image


```
import numpy as np
import mahotas
import mahotas.demos

from mahotas.thresholding import soft_threshold
from matplotlib import pyplot as plt
from os import path
f = mahotas.demos.load('luispedro', as_grey=True)
f = f[:256,:256]
plt.gray()
# Show the data:
print("Fraction of zeros in original image: {0}".format(np.mean(f==0)))
plt.imshow(f)
plt.show()
```



There are no zeros in the original image. We now try a baseline compression method: save every other pixel and only high-order bits.


```
direct = f[::2, ::2].copy()
direct /= 8
direct = direct.astype(np.uint8)
print("Fraction of zeros in original image (after division by 8): {}".format(np.
    mean(direct==0)))
plt.imshow(direct)
plt.show()
```



There are only a few zeros, though. We have, however, thrown away 75% of the values. Can we get a better image, using the same number of values, though?


We will transform the image using a Daubechies wavelet (D8) and then discard the high-order bits.

```
# Transform using D8 Wavelet to obtain transformed image t:  
t = mahotas.daubechies(f, 'D8')  
  
# Discard low-order bits:  
t /= 8  
t = t.astype(np.int8)  
print("Fraction of zeros in transform (after division by 8): {}".format(np.  
    mean(t==0)))  
plt.imshow(t)  
plt.show()
```



This has 60% zeros! What does the reconstructed image look like?


```
# Let us look at what this looks like
r = mahotas.idaubechies(t, 'D8')
plt.imshow(r)
plt.show()
```



This is a pretty good reduction without much quality loss. We can go further and discard small values in the transformed space. Also, let's make the remaining values even smaller in magnitude.

Now, this will be 77% of zeros, with the remaining being small values. This image would compress very well as a lossless image and we could reconstruct the full image after transmission. The quality is certainly higher than just keeping every fourth pixel and low-order bits.

```
tt = soft_threshold(t, 12)
print("Fraction of zeros in transform (after division by 8 & soft thresholding): {0}".
      format(np.mean(tt==0)))
# Let us look again at what we have:
rt = mahotas.idaubechies(tt, 'D8')
plt.imshow(rt)
```



What About the Borders?

In this example, we can see some artifacts at the border. We can use `wavelet_center` and `wavelet_decenter` to handle borders to correctly:

```
fc = mahotas.wavelet_center(f)
t = mahotas.daubechies(fc, 'D8')
r = mahotas.idaubechies(fc, 'D8')
rd = mahotas.wavelet_decenter(r, fc.shape)
```

Now, `rd` is equal (except for rounding) to `fc` **without any border effects**.

API Documentation

A package for computer vision in Python.

Main Features

features Compute global and local features (several submodules, include SURF and Haralick features)

convolve Convolution and wavelets

morph Morphological features. Most are available at the `mahotas` level, include `erode()`, `dilate()`...

watershed Seeded watershed implementation

imread/imsave read/write image

Documentation: <https://mahotas.readthedocs.io/>

Citation:

Coelho, Luis Pedro, 2013. Mahotas: Open source software for scriptable computer vision. Journal of Open Research Software, 1:e3, DOI: <http://dx.doi.org/10.5334/jors.ac>

`mahotas.haar(f, preserve_energy=True, inline=False)`

Haar transform

Parameters `f` : 2-D ndarray

Input image

`preserve_energy` : bool, optional

Whether to normalise the result so that energy is preserved (the default).

`inline` : bool, optional

Whether to write the results to the input image. By default, a new image is returned.

Integer images are always converted to floating point and copied.

See also:

`ihaar` function Reverse Haar transform

`mahotas.ihaar(f, preserve_energy=True, inline=False)`

Reverse Haar transform

`ihaar(haar(f))` is more or less equal to `f` (equal, except for possible rounding issues).

Parameters `f` : 2-D ndarray

Input image. If it is an integer image, it is converted to floating point (double).

`preserve_energy` : bool, optional

Whether to normalise the result so that energy is preserved (the default).

`inline` : bool, optional

Whether to write the results to the input image. By default, a new image is returned.

Integer images are always converted to floating point and copied.

Returns `f` : ndarray

See also:

`haar` function Forward Haar transform

`mahotas.daubechies(f, code, inline=False)`

Daubechies wavelet transform

This function works best if the image sizes are powers of 2!

Parameters `f` : ndarray

2-D image

`code` : str

One of ‘D2’, ‘D4’, ... ‘D20’

inline : bool, optional

Whether to write the results to the input image. By default, a new image is returned. Integer images are always converted to floating point and copied.

See also:

`haar` function Haar transform (equivalent to D2)

`mahotas.idaubechies(f, code, inline=False)`

Daubechies wavelet inverse transform

Parameters `f` : ndarray

2-D image

code : str

One of ‘D2’, ‘D4’, ... ‘D20’

inline : bool, optional

Whether to write the results to the input image. By default, a new image is returned. Integer images are always converted to floating point and copied.

See also:

`haar` function Haar transform (equivalent to D2)

Distance Transform

The example in this section is present in the source under `mahotas/demos/distance.py`.

We start with an image, a black&white image that is mostly black except for two white spots:


```
import numpy as np
import mahotas

f = np.ones((256,256), bool)
f[200:,240:] = False
f[128:144,32:48] = False
```

```
from pylab import imshow, gray, show
import numpy as np

f = np.ones((256,256), bool)
f[200:,240:] = False
f[128:144,32:48] = False

gray()
imshow(f)
show()
```



There is a simple `distance()` function which computes the distance map:

```
import mahotas
dmap = mahotas.distance(f)
```


Now `dmap[y, x]` contains the squared euclidean distance of the pixel (y,x) to the nearest black pixel in `f`. If `f[y, x] == True`, then `dmap[y, x] == 0`.

```
from __future__ import print_function

import pylab as p
import numpy as np
import mahotas

f = np.ones((256,256), bool)
f[200:,240:] = False
f[128:144,32:48] = False
# f is basically True with the exception of two islands: one in the lower-right
# corner, another, middle-left

dmap = mahotas.distance(f)
p.imshow(dmap)
p.show()
```



Distance Transform and Watershed

The distance transform is often combined with the watershed for segmentation. Here is an example (which is available with the source in the `mahotas/demos/` directory as `nuclear_distance_watershed.py`).

```

import mahotas
from os import path
import numpy as np
from matplotlib import pyplot as plt


try:
    nuclear_path = path.join(
        path.dirname(path.abspath(__file__)),
        'data',
        'nuclear.png')
except NameError:
    nuclear_path = path.join('data', 'nuclear.png')

nuclear = mahotas.imread(nuclear_path)
nuclear = nuclear[:, :, 0]
nuclear = mahotas.gaussian_filter(nuclear, 1.)
thresed = (nuclear > nuclear.mean())
distances = mahotas.stretch(mahotas.distance(thresed))
Bc = np.ones((9, 9))

```

```
maxima = mahotas.morph.regmax(distances, Bc=Bc)
spots,n_spots = mahotas.label(maxima, Bc=Bc)
surface = (distances.max() - distances)
areas = mahotas.cwatershed(surface, spots)
areas *= threshed

import random
from matplotlib import colors as c
colors = map(plt.cm.jet, range(0, 256, 4))
random.shuffle(colors)
colors[0] = (0.,0.,0.,1.)
rmap = c.ListedColormap(colors)
plt.imshow(areas, cmap=rmap)
plt.show()
```



The code is not very complex. Start by loading the image and preprocessing it with a Gaussian blur:

```
import mahotas
import mahotas.demos

nuclear = mahotas.demos.nuclear_image()
nuclear = nuclear[:, :, 0]
```

```
nuclear = mahotas.gaussian_filter(nuclear, 1.)
threshed = (nuclear > nuclear.mean())
```

Now, we compute the distance transform:

```
distancess = mahotas.stretch(mahotas.distance(threshed))
```

We find and label the regional maxima:

```
Bc = np.ones((9,9))
maxima = mahotas.morph.regmax(distancess, Bc=Bc)
spots,n_spots = mahotas.label(maxima, Bc=Bc)
```

Finally, to obtain the image above, we invert the distance transform (because of the way that `cwatershed` is defined) and compute the watershed:

```
surface = (distancess.max() - distancess)
areas = mahotas.cwatershed(surface, spots)
areas *= threshed
```

We used a random colormap with a black background for the final image. This is achieved by:

```
import random
from matplotlib import colors as c
colors = map(cm.jet, range(0, 256, 4))
random.shuffle(colors)
colors[0] = (0., 0., 0., 1.)
rmap = c.ListedColormap(colors)
imshow(areas, cmap=rmap)
show()
```

API Documentation

A package for computer vision in Python.

Main Features

features Compute global and local features (several submodules, include SURF and Haralick features)

convolve Convolution and wavelets

morph Morphological features. Most are available at the mahotas level, include erode(), dilate()...

watershed Seeded watershed implementation

imread/imsave read/write image

Documentation: <https://mahotas.readthedocs.io/>

Citation:

Coelho, Luis Pedro, 2013. Mahotas: Open source software for scriptable computer vision. Journal of Open Research Software, 1:e3, DOI: <http://dx.doi.org/10.5334/jors.ac>

`mahotas.distance(bw, metric='euclidean2')`

Computes the distance transform of image `bw`:

```
dmap[i, j] = min_{i', j'} { (i-i')**2 + (j-j')**2 | !bw[i', j'] }
```

That is, at each point, compute the distance to the background.

If there is no background, then a very high value will be returned in all pixels (this is a sort of infinity).

Parameters `bw` : ndarray

If boolean, `False` will denote the background and `True` the foreground. If not boolean, this will be interpreted as `bw != 0` (this way you can use labeled images without any problems).

metric : str, optional

one of ‘euclidean2’ (default) or ‘euclidean’

Returns `dmap` : ndarray

distance map

References

For 2-D images, the following algorithm is used:

Felzenszwalb P, Huttenlocher D. *Distance transforms of sampled functions*. Cornell Computing and Information. 2004.

Available at: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.88.1647&rep=rep1&type=pdf>.

For n-D images (with $n > 2$), a slower hand-craft method is used.

Polygon Utilities

Drawing

Mahotas is not a package to generate images, but there are a few simple functions to draw lines and polygons on an image (the target image is known as the *canvas* in this documentation).

The simplest function is `line`: Give it two points and it draws a line between them. The implementation is simple, and in Python, so it will be slow for many complex usage.

The main purpose of these utilities is to aid debugging and visualisation. If you need to generate fancy graphs, look for packages such as `matplotlib`.

Convex Hull

Convex hull functions are a more typical image processing feature. Mahotas has a simple one, called `convexhull`. Given a boolean image (or anything that will get interpreted as a boolean image), it finds the convex hull of all its on points.

The implementation is in C++, so it is fast.

A companion function `fill_convexhull` returns the convex hull as a binary image.

API Documentation

`mahotas.polygon.line((y0, x0), (y1, x1), canvas, color=1)`
Draw a line

Parameters `p0` : pair of integers
 first point
`p1` : pair of integers
 second point
`canvas` : ndarray
 where to draw, will be modified in place
`color` : integer, optional
 which value to store on the pixels (default: 1)

Notes

Implementation Reference

`mahotas.polygon.fill_polygon([(y0, x0), (y1, x1), ...], canvas, color=1)`
Draw a filled polygon in canvas

Parameters `polygon` : list of pairs
 a list of (y,x) points
`canvas` : ndarray
 where to draw, will be modified in place
`color` : integer, optional
 which colour to use (default: 1)

`mahotas.polygon.convexhull(bwimg)`

Compute the convex hull as a polygon

This is an implementation of the Graham Scan: http://en.wikipedia.org/wiki/Graham_scan

Parameters `bwimg` : ndarray
 input image (interpreted as boolean). Only 2D arrays are supported.

Returns `hull` : ndarray
 Set of (y,x) coordinates of hull corners

`mahotas.polygon.fill_convexhull(bwimg)`

Compute the convex hull and return it as a binary mask

Parameters `bwimage` : input image (interpreted as boolean)
Returns `hull` : image of same size and dtype as `bwimg` with the hull filled in.

Features

By features we mean, basically, numerical functions of the image. That is, any method that gives me a number from the image, I can call it a *feature*. Ideally, these should be meaningful.

We can classify features into two types:

global These are a function of the whole image.

local These **have a position** and are a function of a local image region.

Mahotas supports both types.

The classification tutorial illustrates the usefulness of feature computation.

If you simply want to compute features from images (without any further processing), you can also use the `mahotas-features.py` script, which is installed with mahotas (since version 1.4).

Global features

Haralick features

These are texture features, based on the adjacency matrix (the adjacency matrix stores in position (i,j) the number of times that a pixel takes the value i **next to** a pixel with the value j). Given different ways to define **next to**, you obtain slightly different variations of the features. Standard practice is to average them out across the directions to get some rotational invariance.

They can be computed for 2-D or 3-D images and are available in the `mahotas.features.haralick` module.

Only the first 13 features are implemented. The last (14th) feature is normally considered to be *unstable*, although it is not clear to me why this is. (See this unanswered question on [Cross-validated](#)).

Local Binary Patterns

Local binary patterns (LBP) are a more recent set of features. Each pixel is looked at individually. Its neighbourhood is analysed and summarised by a single numeric code. The normalised histogram across all the pixels in the image is the final set of features.

Again, this is an attempt at capturing texture. LBPs are insensitive to orientation and to illumination (scaling).

Threshold Adjacency Statistics

Threshold adjacency statistics (TAS) are a recent innovation too. In the original version, they have fixed parameters, but we have adapted them to *parameter-free* versions (see [Structured Literature Image Finder: Extracting Information from Text and Images in Biomedical Literature](#) by Coelho et al. for a reference). Mahotas supports both.

Zernike Moments

Zernike moments are **not** a texture feature, but rather a global measure of how the mass is distributed.

Local features

SURF: Speeded-Up Robust Features

Speeded-Up Robust Features (SURF) have both a *location* (pixel coordinates) and a *scale* (natural size) as well as a descriptor (the local features).

Read more about SURF.

Local Binary Patterns

Local binary patterns depend on the local region around each pixel. See the diagram below:

(Image reference: [Wikipedia](#))

The reference pixel is in red, at the centre. A number of points are defined at a distance r from it. These are the green points. As you go from left to right, the number of green points increases.

The “pattern” in the name is the relationship of the value at the green points when compared to the central red point. We call it a binary pattern because all that is taken into account is whether the value at the green point is greater than the value at the red point.

As you can see, the green points do not necessarily fall exactly on another pixel, so we need to use interpolation to find a value for the green points.

API Documentation

The `mahotas.features.lb` module contains the `lbp` function which implements LBPs.

`mahotas.features.lbp(image, radius, points, ignore_zeros=False)`

Compute Linear Binary Patterns

The return value is a **histogram** of feature counts, where position i corresponds to the number of pixels that had code i . The codes are compressed so that impossible codes are not used. Therefore, this is the i^{th} feature, not just the feature with binary code `` i ''.

Parameters `image` : ndarray

 input image (2-D numpy ndarray)

`radius` : number (integer or floating point)

 radius (in pixels)

`points` : integer

 nr of points to consider

`ignore_zeros` : boolean, optional

 whether to ignore zeros (default: False)

Returns `features` : 1-D numpy ndarray

 histogram of features. See above for a caveat on the interpretation of these.

References

Gray Scale and Rotation Invariant Texture Classification with Local Binary Patterns Ojala, T. Pietikainen, M. Maenpaa, T. Lecture Notes in Computer Science (Springer) 2000, ISSU 1842, pages 404-420

```
mahotas.features.lbp.lbp_transform(image, radius, points, ignore_zeros=False, preserve_shape=True)
```

Compute Linear Binary Pattern Transform

The return value are the transformed pixel values **histogram** of feature counts, where position *i* corresponds to the number of pixels that had code *i*. The codes are compressed so that impossible codes are not used. Therefore, this is the *i`th feature*, not just the feature with binary code ``*i*''.

Parameters **image** : ndarray

input image (2-D numpy ndarray)

radius : number (integer or floating point)

radius (in pixels)

points : integer

nr of points to consider

ignore_zeros : boolean, optional

whether to ignore zeros. Note that if you set this to True, you will need to set preserve_shape to False. (default: False)

preserve_shape : boolean, optional

whether to return an array with the same shape as **image**. (default: True)

Returns **features** : 1-D numpy ndarray

histogram of features. See above for a caveat on the interpretation of these.

References

Gray Scale and Rotation Invariant Texture Classification with Local Binary Patterns Ojala, T. Pietikainen, M. Maenpaa, T. Lecture Notes in Computer Science (Springer) 2000, ISSU 1842, pages 404-420

Speeded-Up Robust Features

New in version 0.8: In version 0.8, some of the inner functions are now in `mahotas.features.surf` instead of `mahotas.surf`. Speeded-Up Robust Features (**SURF**) are a recent innovation in the *local features* family. There are two steps to this algorithm:

1. Detection of interest points.
2. Description of interest points.

The function `mahotas.features.surf.surf` combines the two steps:

```
import numpy as np
from mahotas.features import surf

f = ... # input image
spoints = surf.surf(f)
print("Nr points: {}".format(len(spoints)))
```

Given the results, we can perform a simple clustering using, for example, `milk` (nowadays, scikit-learn would be a better choice):

```
try:
    import milk

    # spoints includes both the detection information (such as the position
    # and the scale) as well as the descriptor (i.e., what the area around
    # the point looks like). We only want to use the descriptor for
    # clustering. The descriptor starts at position 5:
    descrs = spoints[:,5:]

    # We use 5 colours just because if it was much larger, then the colours
    # would look too similar in the output.
    k = 5
    values, _ = milk.kmeans(descrs, k)
    colors = np.array([(255-52*i, 25+52*i, 37**i % 101) for i in xrange(k)])
except:
    values = np.zeros(100)
    colors = [(255,0,0)]
```

So we are assigning different colours to each of the possible

The helper `surf.show_surf` draws coloured polygons around the interest points:

```
f2 = surf.show_surf(f, spoints[:100], values, colors)
imshow(f2)
show()
```

Running the above on a photo of luispedro, the author of mahotas yields:

```
from __future__ import print_function
import numpy as np
import mahotas as mh
from mahotas.features import surf
from pylab import *


from os import path

f = mh.demos.load('luispedro', as_grey=True)
f = f.astype(np.uint8)
spoints = surf.surf(f, 4, 6, 2)
print("Nr points:", len(spoints))

try:
    import milk
    descrs = spoints[:,5:]
    k = 5
    values, _ = milk.kmeans(descrs, k)
    colors = np.array([(255-52*i, 25+52*i, 37**i % 101) for i in range(k)])
except:
```

```
values = np.zeros(100)
colors = np.array([(255, 0, 0)])

f2 = surf.show_surf(f, spoints[:100], values, colors)
imshow(f2)
show()
```



API Documentation

The `mahotas.features.surf` module contains separate functions for all the steps in the SURF pipeline.

`mahotas.features.surf.dense(f, spacing, scale={np.sqrt(spacing)}, is_integral=False, include_interest_point=False)`

Parameters `f` : image

original image

`spacing` : integer

Distance between points

`scale` : float, optional

Scale of interest points. By default, it is set to `np.sqrt(spacing)`

`is_integral` : boolean, optional

Whether f is an integral image

include_interest_point : bool, optional

Whether to return interest point information. Default is False

Returns descriptors : ndarray

Descriptors at dense points. Note that the interest point is **not returned by default**.

See also:

surf function Find interest points and then compute descriptors

descriptors function Compute descriptors at user provided interest points

```
mahotas.features.surf.integral(f, in_place=False, dtype=<type 'numpy.float64'>)
fi = integral(f, in_place=False, dtype=np.double):
```

Compute integral image

Parameters f : ndarray

input image. Only 2-D images are supported.

in_place : bool, optional

Whether to overwrite f (default: False).

dtype : dtype, optional

dtype to use (default: double)

Returns fi : ndarray of *dtype* of same shape as f

The integral image

```
mahotas.features.surf.surf(f, nr_octaves=4, nr_scales=6, initial_step_size=1, threshold=0.1,
                           max_points=1024, descriptor_only=False)
points = surf(f, nr_octaves=4, nr_scales=6, initial_step_size=1, threshold=0.1, max_points=1024, descriptor_only=False):
```

Run SURF detection and descriptor computations

Speeded-Up Robust Features (SURF) are fast local features computed at automatically determined keypoints.

Parameters f : ndarray

input image

nr_octaves : integer, optional

Nr of octaves (default: 4)

nr_scales : integer, optional

Nr of scales (default: 6)

initial_step_size : integer, optional

Initial step size in pixels (default: 1)

threshold : float, optional

Threshold of the strength of the interest point (default: 0.1)

max_points : integer, optional

Maximum number of points to return. By default, return at most 1024 points. Note that the number may be smaller even in the case where there are that many points. This is a side-effect of the way the threshold is implemented: only `max_points` are considered, but some of those may be filtered out.

descriptor_only : boolean, optional

If `descriptor_only`, then returns only the 64-element descriptors (default is `False`).

Returns points : ndarray of double, shape = (N, 6 + 64)

N is nr of points. Each point is represented as (*y,x,scale,score,laplacian,angle,D_0,...,D_63*) where *y,x,scale* is the position, *angle* the orientation, *score* and *laplacian* the score and sign of the detector; and *D_i* is the descriptor

If `descriptor_only`, then only the *D_is* are returned and the array has shape (N, 64)!

References

Herbert Bay, Andreas Ess, Tinne Tuytelaars, Luc Van Gool “SURF: Speeded Up Robust Features”, Computer Vision and Image Understanding (CVIU), Vol. 110, No. 3, pp. 346–359, 2008

Implementing SURF-ref With Mahotas

This is a companion to the paper Determining the subcellular location of new proteins from microscope images using local features by Coelho et al. (2013).

```
def surf_ref(f, ref):
    """
    features = surf_ref(f, ref)

    Computer SURF-ref features

    Parameters
    -----
    f : ndarray
        input image
    ref : ndarray
        Corresponding reference image

    Returns
    -----
    features : ndarray
        descriptors
    """

    fi = surf.integral(f.copy())
    points = surf.interest_points(fi, 6, 24, 1, max_points=1024, is_integral=True)
    descs = surf.descriptors(fi, points, is_integral=True, descriptor_only=True)
    if ref is None:
        return descs
    descsref = surf.descriptors(ref, points, descriptor_only=True)
    return np.hstack( (descs, descsref) )
```

This function can take any number of reference images.

We now compute all features for all images in widefield dataset:

```
from glob import glob
import re

basedir = 'rt-widefield' # Edit as needed

features = []
labels = []

# We need the following to keep track of the proteins:
origins = []
prev_origin = ''
origin_counter = -1 # set to -1 so it will be correctly initialized on the first image

for dir in glob(basedir):
    if dir == 'README': continue
    for f in glob('{}/{}/*-protein.tif'.format(basedir, dir)):
        origin = f[:5]
        if origin != prev_origin:
            origin_counter += 1
            prev_origin = origin

        f = '{}/{}/{}'.format(basedir, dir, f)
        f = mh.imread(f)
        ref = mh.imread(f.replace('protein', 'dna'))
        features.append(surf_ref(f, ref))
        labels.append(dir)
        origins.append(origin_counter)
```

Classification

With all the precomputed features, we can now run 10-fold cross-validation on these features.

We will use milk for machine learning:

```
import milk
```

Milk's interface is around learner objects. We are going to define a function:

```
def train_model(features, labels):
```

The first step is to find centroids:

```
# concatenate all the features:
concatenated = np.concatenate(features)
```

We could use the whole array concatenated for kmeans. However, that would take a long time, so we will use just 1/16th of it:

```
concatenated = concatenated[::-16]
_, centroids = milk.kmeans(concatenated, k=len(labels)//4, R=123)
```

The R argument is the random seed. We set it to a constant to get reproducible results, but feel free to vary it.

Based on these centroids, we project the features to histograms. Now, we are using all of the features:

```
features = np.array([
    project_centroids(centroids, fs, histogram=True)
    for fs in features])
```

Finally, we can use a traditional milk learner (which will perform feature selection, normalization, and SVM training):

```
learner = milk.defaultlearner()
model = learner.train(features, labels)
```

We must return both the centroids that were used and the classification model:

```
return centroids, model
```

To classify an instance, we define another function, which uses the centroids and the model:

```
def apply_many(centroids, model, features):
    features = np.array([
        project_centroids(centroids, fs, histogram=True)
        for fs in features])
    return model.apply_many(features)
```

In fact, while the above will work well, milk already provides a learner object which will perform all of those tasks!

```
import milk
from milk.supervised.precluster import frac_precluster_learner

learner = frac_precluster_learner(kfrac=4, sample=16)
cmatrix, names = milk.nfoldcrossvalidation(features, labels, origins=origins,
                                             learner=learner)
acc = cmatrix.astype(float).trace() / cmatrix.sum()
print('Accuracy: {:.1}%'.format(100.*acc))
```

Morphological Operators


New in version 0.8: open() & close() were added in version 0.8

Morphological operators were the first operations in mahotas (back then, it was even, briefly, just a single C++ module called `morph`). Since then, mahotas has grown a lot. This module, too, has grown and acquired more morphological operators as well as being optimised for speed.

Let us first select an interesting image

```
import mahotas
import mahotas.demos
from pylab import gray, imshow, show
import numpy as np

luispedro = mahotas.demos.load('luispedro')
luispedro = luispedro.max(2)
T = mahotas.otsu(luispedro)
lpbin = (luispedro > T)
gray()
eye = ~lpbin[112:180, 100:190]
imshow(eye)
show()
```



After Oct 2013, you can get this image with mahotas as:

```
import mahotas.demos
luispedro = mahotas.demos.load('luispedro')
luispedro = luispedro.max(2)
```

Dilation & Erosion


Dilation and erosion are two very basic operators (mathematically, you only need one of them as you can define the erosion as dilation of the negative or vice-versa).

These operations are available in the `mahotas.morph` module:

```
mahotas.morph.dilate(eye)
```

Dilation is, intuitively, making positive areas “fatter”:


```
imshow(mahotas.morph.dilate(eye))
show()
```



```
mahotas.morph.erode(eye)
```

Erosion, by contrast, thins them out:

```
imshow(mahotas.morph.erode(eye))  
show()
```



Mahotas supports greyscale erosion and dilation (depending on the `dtype` of the arguments) and you can specify any structuring element you wish (including non-flat ones). By default, a 1-cross is used:

```
# if no structure-element is passed, use a cross:
se = np.array([
    [0, 1, 0],
    [1, 1, 1],
    [0, 1, 0]], bool)
```

However, you can use whatever structuring element you want:

```
se = np.array([
    [1, 1, 0],
    [1, 1, 1],
    [0, 1, 1]], bool)
dilated = mahotas.morph.dilate(eye, se)
eroded = mahotas.morph.erode(eye, se)
```

Note that when you pass it a non-boolean array as the first argument, you will get *greyscale erosion*. Mahotas supports full greyscale erosion, including arbitrary, flat or non-flat, structuring elements).


Close & Open

Closing and opening are based on erosion and dilation. Again, they work in greyscale and can use an arbitrary structure element.

Here is closing:

```
mahotas.morph.close(eye)
```


```
imshow(mahotas.morph.close(eye))  
show()
```



And here is opening:

```
mahotas.morph.open(eye)
```

```
imshow(mahotas.morph.open(eye))  
show()
```



Both `close` and `open` take an optional structuring element as a second argument:

```
mahotas.morph.open(eye, se)
```

Color Space Conversions

New in version 0.9.6.

Red-green-blue images

An RGB image is represented as a 3-dimensional array of shape `(h, w, 3)`, where each pixel is represented by three values, red/green/blue.

For example, the classic lena image is a `(512, 512, 3)` array:


```
import mahotas as mh
lena = mh.demos.load('lena')
print(lena.shape)
```

We can convert it to greyscale as using `rgb2grey` (or `rgb2gray` if you prefer, both work). This conversion uses a visually realistic method (which weighs the green channel more heavily as human eyes are more sensitive to it). For example:

```
import mahotas as mh
lena = mh.demos.load('lena')
lenag = mh.colors.rgb2grey(lena)
```

```
from pylab import imshow
import mahotas as mh
lena = mh.demos.load('lena')
lenag = mh.colors.rgb2grey(lena)

imshow(lenag)
```




We can also convert to sepia with `rgb2sepia`:

```
lenas = mh.colors.rgb2sepia(lena)
```

```
from pylab import imshow
import mahotas as mh
lena = mh.demos.load('lena')

lenas = mh.colors.rgb2sepia(lena)

imshow(lenas)
```



Other Colour Spaces

Mahotas can also convert to [XYZ space](#) and to the [Lab space](#) with `rgb2xyz` and `rgb2lab`, respectively.

Input/Output with Mahotas

Mahotas does not have any builtin support for input/output. However, it wraps a few other libraries that do. The result is that you can do:

```
import mahotas as mh
image = mh.imread('file.png')
mh.imsave('copy.png', image)
```

It can use the following backends (it tries them in the following order):

1. It prefers `mahotas-imread`, if it is available. Imread is a native C++ library which reads images into Numpy arrays. It supports PNG, JPEG, TIFF, WEBP, BMP, and a few TIFF-based microscopy formats (LSM and STK).
2. It also looks for `freeimage`. Freeimage can read and write many formats. Unfortunately, it is harder to install and it is not as well-maintained as `imread`.
3. Finally, it tries to load `pillow`.

Thus, to use the `imread` or `imsave` functions, you need to install one of the packages above. At one point, mahotas supported wrapping matplotlib, but their image loading methods are unreliable as it uses other packages itself. Thus, depending on what you had installed, the resulting images would be different.

If you are running on Windows, you may wish to try [Christoph Gohlke's packages](#).

Tutorial: Classification Using Mahotas

New in version 0.8: Before version 0.8, texture was under `mahotas`, not under `mahotas.features`

Here is an example of using mahotas and `milk` for image classification (but most of the code can easily be adapted to use another machine learning package). I assume that there are three important directories: `positives/` and `negatives/` contain the manually labeled examples, and the rest of the data is in an `unlabeled/` directory.

Here is the simple algorithm:

1. Compute features for all of the images in `positives` and `negatives`
2. learn a classifier
3. use that classifier on the unlabeled images

In the code below I used `jug` to give you the possibility of running it on multiple processors, but the code also works if you remove every line which mentions `TaskGenerator`.

We start with a bunch of imports:

```
from glob import glob
import mahotas
import mahotas.features
import milk
from jug import TaskGenerator
```

Now, we define a function which computes features. In general, texture features are very fast and give very decent results:

```
@TaskGenerator
def features_for(imname):
    img = mahotas.imread(imname)
    return mahotas.features.haralick(img).mean(0)
```

`mahotas.features.haralick` returns features in 4 directions. We just take the mean (sometimes you use the spread `ptp()` too).

Now a pair of functions to learn a classifier and apply it. These are just `milk` functions:

```
@TaskGenerator
def learn_model(features, labels):
    learner = milk.defaultclassifier()
    return learner.train(features, labels)

@TaskGenerator
def classify(model, features):
    return model.apply(features)
```

We assume we have three pre-prepared directories with the images in jpeg format. This bit you will have to adapt for your own settings:

```
positives = glob('positives/*.jpg')
negatives = glob('negatives/*.jpg')
unlabeled = glob('unlabeled/*.jpg')
```

Finally, the actual computation. Get features for all training data and learn a model:

```
features = map(features_for, negatives + positives)
labels = [0] * len(negatives) + [1] * len(positives)

model = learn_model(features, labels)

labeled = [classify(model, features_for(u)) for u in unlabeled]
```

This uses texture features, which is probably good enough, but you can play with other features in `mahotas.features` if you'd like (or try `mahotas.surf`, but that gets more complicated).

(This was motivated by a question on Stackoverflow).

Tutorial: Extended Depth of Field

This is an example of how to use mahotas to implement an algorithm that it does not have built-in: *extended depth of field*.

The idea is that you have a stack of images, taken at different focal points, and you build a single image so that you get everything in focus.

Start with standard imports:

```
import numpy as np
import mahotas as mh
```

We are going to assume that you have an `image` object, which has three dimensions: the stack, height, and width:

```
stack, h, w = image.shape
```

We use `mh.sobel` as the measure of “infocusness” for each pixel¹:

```
focus = np.array([mh.sobel(t, just_filter=True) for t in image])
```

Now, we select the best slice at each pixel location:

```
best = np.argmax(focus, 0)
```

So far, very easy. The next part is the hard part. We want to do the following:

```
r = np.zeros((h,w))-1
for y in xrange(h):
    for x in xrange(w):
        r[y,x] = image[best[y,x], y, x]
```

But this is very slow (never run nested loops in Python if you can avoid it). We get the same result with a slightly less legible, but faster manipulation²:

¹ Other methods simply use a different measure here.

² I am not 100% convinced that this is the best. After all we create an array of size `len(image)` just to index. I would be happy to find an alternative.


```
image = image.reshape((stack,-1)) # image is now (stack, nr_pixels)
image = image.transpose() # image is now (nr_pixels, stack)
r = image[np.arange(len(image)), best.ravel()] # Select the right pixel at each
    ↪location
r = r.reshape((h,w)) # reshape to get final result
```

Et voilà!

Here is an example, from a stack of microbes imaged. This is the maximum intensity projection:



This is the *most in-focus slice* (using the sobel operator as the measure):



And this is the extended depth of field result:



It is clearly sharper, perhaps at the expense of some possible noise. I actually played around with blurring the image a little bit and it did improve things ever so slightly.

mahotas-features.py

New in version 1.4.0: The `mahotas-features.py` script appeared in version 1.4.0 (July 2015)

With the installation of mahotas, a script called `mahotas-features.py` is installed, which can be used to compute features from a set of files.

Usage

You call the script with a set of flags specifying which features you want to compute, followed by a list of files. For example:

```
$ mahotas-features.py --haralick --lbp image-file1.tiff image-file2.tiff
```

This will output to the file `features.tsv` (this default can be changed with the `--output` option)

Full Usage Information

You can obtain help on all the options by running `mahotas-features.py --help`:

```
If you use mahotas in a scientific publication, please cite  
Coelho, LP (2013). http://dx.doi.org/10.5334/jors.ac
```

```
usage: mahotas-features.py [-h] [--output OUTPUT] [--clobber]  
                           [--convert-to-bw CONVERT_TO_BW] [--no-color]  
                           [--haralick] [--lbp] [--lbp-radius LBP_RADIUS]  
                           [--lbp-points LBP_POINTS]  
                           input_file_name [input_file_name ...]
```

```
Compute features using mahotas
```

```
positional arguments:
  input_file_name      Image files names

optional arguments:
  -h, --help            show this help message and exit
  --output OUTPUT       Output file for feature files
  --clobber             Overwrite output file (if it exists)
  --convert-to-bw CONVERT_TO_BW
                        Convert color images to greyscale. Acceptable values:
                        no: raises an error (default) max: use max projection
                        yes: use rgb2gray
  --no-color            Do not print in color (for error and warning messages)
  --haralick            Compute Haralick features
  --lbp                 Compute LBP (linear binary patterns) features
  --lbp-radius LBP_RADIUS
                        Radius to use for LBP features
  --lbp-points LBP_POINTS
                        Nr of points to use for LBP features
```

Frequently Asked Questions

How do I install mahotas with anaconda?

If you are using `conda`, you can install mahotas from `conda-forge` using the following commands:

```
conda config --add channels conda-forge
conda install mahotas
```

Who uses mahotas?

In June 2016, there were 34 papers citing the mahotas paper

Why did you not simply contribute to `scipy.ndimage` or `scikits.image`?

When I started this project (although it wasn't called mahotas and it was more of a collection of semi-organised routines than a project), there was no `scikits.image`.

In the meanwhile, all these projects have very different internal philosophies. `ndimage` is old-school `scipy`, in C, with macros. `scikits.image` uses Cython extensively, while `mahotas` uses C++ and templates. I don't want to use Cython as I find that it is not yet established enough and at the time it could not be used to write functions that run on multiple types (like with C++ templates). The `scipy` community does not want to use C++.

I have, on the other hand, taken code from `ndimage` and ported it to C++ for use in `mahotas`. In the process, I feel it is much cleaner code (because you can use RAII, exceptions, and templates) and I want to keep it that way.

In any case, we all use the same data format: numpy arrays. It is very easy (trivial, really) to use all the packages together and take whatever functions you want from each. All the packages use function based interfaces which make it easy to mix-and-match.

I ran out of memory computing Haralick features on 16 bit images. Is it not supported?

Yes, it is supported, but your machine may not be big enough to do the computation. In order to compute Haralick features, first a cooccurrence matrix is computed. This matrix has the size (`ngrey`, `ngrey`) where `ngrey` is the largest grey value in the input. Thus, if your image has a very high dynamic range (i.e., `ngrey` is large), you may not have the resources to compute the cooccurrence matrix.

It is often a good idea to contrast stretch your images. For example, using the following code, stretches your images to the 0-255 range:

```
im_stretched = mh.stretch(im)
features = mh.features.haralic(im_stretched)
```

16 bit images where the dynamic range is not too large (for example, some imaging equipment can only really produce 12 bits, so `ngrey < 4096`) are not a problem.

What are the parameters to Local Binary Patterns?

Read the documentation on local binary patterns.

I am using mahotas in a scientific publication, is there a citation?

If you use mahotas in a scientific publication, please cite:

Coelho, L.P. 2013. Mahotas: Open source software for scriptable computer vision. Journal of Open Research Software 1(1), DOI: <http://dx.doi.org/10.5334/jors.ac>

In BibTeX format:

```
@article{coelho:mahotas,
    title = {Mahotas: Open source software for scriptable computer vision},
    author = {Luis Pedro Coelho},
    journal = {Journal of Open Research Software},
    year = {2013},
    volume = {1},
    doi = {10.5334/jors.ac},
    url = {http://dx.doi.org/10.5334/jors.ac}
}
```

This is accessible in Python using `mahotas.citation()`.

Imread cannot find FreeImage

Mahotas itself does not have the functionality to read in images (see the I/O section).

Functions such as `imread` are just a wrapper around one of 2 backends:

1. `mahotas-imread` (i.e., <https://pypi.python.org/pypi/imread>)
2. `FreeImage`

Thus, you need to install one of the packages above. At one point, mahotas supported wrapping matplotlib, but their image loading methods are unreliable as it uses other packages itself. Thus, depending on what you had installed, the resulting images would be different.

If you are running on Windows, you may wish to try [Christoph Gohlke's packages](#).

Mahotas Internals

This section is of interest if you are trying to understand how mahotas works in order to fix something, extend it (patches are always welcome), or use some of its technology in your projects.

Philosophy

Mahotas should not suck.

This is my main development goal and, if I achieve it, this alone should put mahotas in the top ten to one percent of software packages.

Mahotas should have no bugs. None. Ever.

Of course, some creep in. So, we settle for the next best thing: *Mahotas should have no **known bugs***. Whenever a bug is discovered, the top priority is to squash it.

Read the principles of mahotas

C++/Python Division

Mahotas is, for the most part, written in C++, but almost always, you call a Python function which checks types and then calls the internal function. This is slightly slower, but it is easier to develop this way (and, for all but the smallest image, it will not matter).

So each `module.py` will have its associated `_module.cpp`.

C++ Templates

The main reason that mahotas is in C++ (and not in pure C) is to use templates. Almost all C++ functions are actually 2 functions:

1. A `py_function` which uses the Python C/API to get arguments, &c. This is almost always pure C.
2. A template function`<dtype>` which works for the `dtype` performing the actual operation.

So, for example, this is how `erode` is implemented. `py_erode` is generic:

```
PyObject* py_erode(PyObject* self, PyObject* args) {
    PyArrayObject* array;
    PyArrayObject* Bc;
    if (!PyArg_ParseTuple(args, "OO", &array, &Bc)) return NULL;
    PyArrayObject* res_a = (PyArrayObject*)PyArray_SimpleNew(array->nd, array->
    ↪dimensions, PyArray_TYPE(array));
    if (!res_a) return NULL;
    PyArray_FILLWBYTE(res_a, 0);
#define HANDLE(type) \
    erode<type>(numpy::aligned_array<type>(res_a), numpy::aligned_array<type>(array), \
    ↪numpy::aligned_array<type>(Bc)); \
    SAFE_SWITCH_ON_INTEGER_TYPES_OF(array)
#undef HANDLE
    ...
}
```

These functions normally contain a lot of boiler-plate code: read the arguments, perform some sanity checks, perhaps a bit of initialisation, and then, the switch on the input type with the help of the `SAFE_SWITCH_ON_INTEGER_TYPES_OF()` and friends, which call the right specialisation of the template that does the actual work. In this example `erode` implements (binary) erosion:

```
template<typename T>
void erode(numpy::aligned_array<T> res, numpy::aligned_array<T> array, numpy::aligned_
array<T> Bc) {
    gil_release nogil;
    const unsigned N = res.size();
    typename numpy::aligned_array<T>::iterator iter = array.begin();
    filter_iterator<T> filter(res.raw_array(), Bc.raw_array());
    const unsigned N2 = filter.size();
    T* rpos = res.data();

    for (int i = 0; i != N; ++i, ++rpos, filter.iterate_with(iter), ++iter) {
        for (int j = 0; j != N2; ++j) {
            T arr_val = false;
            filter.retrieve(iter, j, arr_val);
            if (filter[j] && !arr_val) goto skip_this_one;
        }
        *rpos = true;
        skip_this_one: continue;
    }
}
```

The template machinery is not that complicated and the functions using it are very simple and easy to read. The only downside is that there is some expansion of code size. Given the small size of these functions however, this is not a big issue.

In the snippet above, you can see some other C++ machinery:

gil_release This is a RAII object that release the GIL in its constructor and gets it back in its destructor. Normally, the template function will release the GIL after the Python-specific code is done.

array This is a thin wrapper around `PyArrayObject` that knows its type and has iterators. Relying on these objects has the further advantage that in debug mode, it checks bounds for many memory accesses. While this is very costly for everyday usage, it can catch bugs faster than the alternatives.

filter_iterator This is taken from `scipy.ndimage` and it is useful to iterate over an image and use a centered filter around each pixel (it keeps track of all of the boundary conditions).

The inner loop is as direct an implementation of erosion as one would wish for: for each pixel in the image, look at its neighbours. If all are true, then set the corresponding output pixel to `true` (else, skip it as it has been initialised to zero).

Most of the functions follow this architecture.

The Why of mahotas

Principles of Mahotas

Here are the principles of mahotas, in decreasing order of importance:

1. Just work
2. Well documented

3. Fast code
4. Simple code
5. Minimal dependencies

Just work

The first principle is that things should *just work*. This means two things: (1) there should be no bugs, and (2) interfaces should be flexible or fail well.

To avoid bugs, tests are extensively used. Every reported bug leads to a new test case, so that it never happens again. New features should at least have a smoke test (test that runs the feature and verifies some basic properties of the output).

Interfaces are designed to be as flexible as possible. No specific types are required unless it is really needed or in performance-enhancing features (such as using `out` parameters).

The user should never be able to crash the Python interpreter with mahotas.

Well documented

No public function is without a complete docstring. In addition to that *hard documentation* (i.e., information with complete technical detail of every nook and cranny of the interface), there is also *soft documentation* (tutorial-like documentation with examples and higher level reasoning).

Fast code

Performance is a feature.

The code should be as fast as possible without sacrificing generality (see *just work* above). This is why C++ templates are used for type independent code.

Simple code

The code should be simple.

Minimal dependencies

Mahotas tries to avoid dependencies.

Right now, building mahotas depends on a C++ compiler, numpy. These are unlikely to ever change. To run mahotas, we need numpy. In order to read images, we need one of (1) imread or (2) FreeImage.

The imread/freeimage dependency is a soft dependency: everything, except for imread works without it. The code is written to ensure that `import-ing` mahotas without an IO backend will not trigger an error unless the `imread()` function is used.

Therefore, once mahotas is compiled, all you really need is numpy. This is unlikely to ever change.

Contributing

Development happens on [github](#) and the preferred contribution method is by forking the repository and issuing a pull request. Alternatively, just sending a patch to luis@luispedro.org will work just as well.

If you don't know git (or another distributed version control system), which is a fantastic tool in general, there are several good git and github tutorials. You can start with their [official documentation](#).

If you want to work on the C++ code, you can read the chapter in the internals before you start. Also, read the principles declaration.

Debug Mode

If you compile mahotas in debug mode, then it will run slower but perform a lot of runtime checks. This is controlled by the `DEBUG` environment variable.

There are two levels:

1. `DEBUG=1` This turns on assertions. The code will run slower, but probably not noticeably slower, except for very large images.
2. `DEBUG=2` This turns on the assertions and additionally uses the debug version of the C++ library (this only works if you are using GCC). Some of the internal code also picks up on the `DEBUG=2` and adds even more sanity checking. The result will be code that runs **much slower** as all operations done through iterators into standard containers are now checked (including many inner loop operations). However, it catches many errors.

The Makefile that comes with the source helps you:

```
make clean  
make debug  
make test
```

will rebuild in debug mode and run all tests. When you are done testing, use the `fast` Make target to get the non-debug build:

```
make clean  
make fast
```

Using `make` will not change your environment. The `DEBUG` variable is set internally only.

If you don't know about it, check out `ccache` which is a great tool if you are developing in compiled languages (this is not specific to mahotas or even Python). It will allow you to quickly perform `make clean`; `make debug` and `make clean`; `make fast` so you never get your builds mixed up.

Possible Tasks

Here are a few ideas for improving mahotas.

New Features

- HOG
- BRISK
- Canny edge detection

- Hough Transform
- bilateral filtering
- Non Local Filtering
- Wiener filtering

Small Improvements

- something like the `overlay` function from `pymorph` (or even just copy it over and adapt it to mahotas style).
- H-maxima transform (again, `pymorph` can provide a basis)
- `entropy thresholding`

Internals

These can be very complex as they require an understanding of the inner workings of mahotas, but that does appeal to a certain personality.

- special case 1-D convolution on C-Arrays in C++. The idea is that you can write a tight inner loop in one dimension:

```
void multiply(floating* r, const floating* f, const floating a, const int n,  
             const int r_step, const int f_step) {  
    for (int i = 0; i != n; ++i) {  
        *r += a * *f;  
        r += r_step;  
        f += f_step;  
    }  
}
```

to implement:

```
r[row] += a* f[row+offset]
```

and you can call this with all the different values of `a` and `offset` that make up your filter. This would be useful for Gaussian filtering.

Tutorials

Mahotas has very good API documentation, but not so many *start to finish* tutorials which touch several parts of it (and even other packages, the ability to seamlessly use other packages in Python is, of course, a good reason to use it).

History

Version 1.4.3 (Oct 3 2016)

- Fix distribution (add missing README.md file)

Version 1.4.2 (Oct 2 2016)

- Fix `resize__to` return exactly the requested size
- Fix hard crash when computing texture on arrays with negative values (issue #72)
- Added `distance` argument to haralick features (pull request #76, by Guillaume Lemaitre)

Version 1.4.1 (Dec 20 2015)

- Add `filter_labeled` function
- Fix tests on 32 bit platforms and older versions of numpy

Version 1.4.0 (July 8 2015)

- Added `mahotas-features.py` script
- Add `short` argument to `citation()` function
- Add `max_iter` argument to `thin()` function
- Fixed `labeled.bbox` when there is no background (issue #61, reported by Daniel Haehn)
- `bbox` now allows dimensions greater than 2 (including when using the `as_slice` and `border` arguments)
- Extended `croptobbox` for dimensions greater than 2
- Added `use_x_minus_y_variance` option to haralick features
- Add function `lbp_names`

Version 1.3.0 (April 28 2015)

- Improve memory handling in `freeimage.write_multipage`
- Fix moments parameter swap
- Add `labeled.bbox` function
- Add `return_mean` and `return_mean_ptp` arguments to haralick function
- Add difference of Gaussians filter (by Jianyu Wang)
- Add Laplacian filter (by Jianyu Wang)
- Fix crash in `median_filter` when mismatched arguments are passed
- Fix `gaussian_filter1d` for `ndim > 2`

Version 1.2.4 (December 23 2014)

- Add PIL based IO

Version 1.2.3 (November 8 2014)

- Export mean_filter at top level
- Fix to Zernike moments computation (reported by Sergey Demurin)
- Fix compilation in platforms without npy_float128 (patch by Gabi Davar)

Version 1.2.2 (October 19 2014)

- Add minlength argument to labeled_sum
- Generalize regmax/regmin to work with floating point images
- Allow floating point inputs to cwatershed()
- Correctly check for float16 & float128 inputs
- Make sobel into a pure function (i.e., do not normalize its input)
- Fix sobel filtering

Version 1.2.1 (July 21 2014)

- Explicitly set numpy.include_dirs() in setup.py [patch by Andrew Stromnov]

Version 1.2 (July 17 2014)

- Export locmaxllocmin at the mahotas namespace level
- Break away ellipse_axes from eccentricity code as it can be useful on its own
- Add find() function
- Add mean_filter() function
- Fix cwatershed() overflow possibility
- Make labeled functions more flexible in accepting more types
- Fix crash in close_holes() with nD images (for n > 2)
- Remove matplotlibwrap
- Use standard setuptools for building (instead of numpy.distutils)
- Add overlay() function

Version 1.1.1 (July 4 2014)

- Fix crash in close_holes() with nD images (for n > 2)

1.1.0 (February 12 2014)

- Better error checking
- Fix interpolation of integer images using order 1
- Add resize_to & resize_rgb_to
- Add coveralls coverage
- Fix SLIC superpixels connectivity
- Add remove_regions_where function
- Fix hard crash in convolution
- Fix axis handling in convolve1d
- Add normalization to moments calculation

1.0.4 (2013-12-15)

- Add mahotas.demos.load()
- Add stretch_rgb() function
- Add demos to mahotas namespace
- Fix SLIC superpixels

1.0.3 (2013-10-06)

- Add border & as_slice arguments to bbox()
- Better error message in gaussian_filter
- Allow as_rgb() to take integer arguments
- Extend distance() to n-dimensions
- Update to newer Numpy APIs (remove direct access to PyArray members)

1.0.2 (July 10 2013)

- Fix requirements filename

1.0.1 (July 9 2013)

- Add lbp_transform() function
- Add rgb2sepia function
- Add mahotas.demos.nuclear_image() function
- Work around matplotlib.imsave's implementation of greyscale
- Fix Haralick bug (report & patch by Tony S Yu)
- Add count_binary1s() function

1.0 (May 21 2013)

- Fix a few corner cases in texture analysis
- Integrate with travis
- Update citation (include DOI)

0.99 (May 4 2013)

- Make matplotlib a soft dependency
- Add demos.image_path() function
- Add citation() function

This version is **1.0 beta**.

0.9.8 (April 22 2013)

- Use matplotlib as IO backend (fallback only)
- Compute dense SURF features
- Fix sobel edge filtering (post-processing)
- Faster 1D convolutions (including faster Gaussian filtering)
- Location independent tests (run mahotas.tests.run() anywhere)
- Add labeled.is_same_labeling function
- Post filter SLIC for smoother regions
- Fix compilation warnings on several platforms

0.9.7 (February 03 2013)

- Add haralick_features function
- Add out parameter to morph functions which were missing it
- Fix erode() & dilate() with empty structuring elements
- Special case binary erosion/dilation in C-Arrays
- Fix long-standing warning in TAS on zero inputs
- Add verbose argument to tests.run()
- Add circle_se to morph
- Allow loc(max|min) to take floating point inputs
- Add Bernsen local thresholding (bernsen and gbernsen functions)

0.9.6 (December 02 2012)

- Fix `distance()` of non-boolean images (issue #24 on github)
- Fix encoding issue on PY3 on Mac OS (issue #25 on github)
- Add `relabel()` function
- Add `remove_regions()` function in labeled module
- Fix `median_filter()` on the borders (respect the `mode` argument)
- Add `mahotas.color` module for conversion between colour spaces
- Add SLIC Superpixels
- Many improvements to the documentation

0.9.5 (November 05 2012)

- Fix compilation in older G++
- Faster Otsu thresholding
- Python 3 support without 2to3
- Add `cdilate` function
- Add `subm` function
- Add tophat transforms (functions `tophat_close` and `tophat_open`)
- Add `mode` argument to `euler()` (patch by Karol M. Langner)
- Add `mode` argument to `bwperim()` & `borders()` (patch by Karol M. Langner)

0.9.4 (October 10 2012)

- Fix compilation on 32-bit machines (Patch by Christoph Gohlke)

0.9.3 (October 9 2012)

- Fix interpolation (Report by Christoph Gohlke)
- Fix second interpolation bug (Report and patch by Christoph Gohlke)
- Update tests to newer numpy
- Enhanced debug mode (compile with `DEBUG=2` in environment)
- Faster `morph.dilate()`
- Add `labeled.labeled_max` & `labeled.labeled_min` (This also led to a refactoring of the `labeled_*` code)
- Many documentation fixes

0.9.2 (September 1 2012)

- Fix compilation on Mac OS X 10.8 (reported by Davide Cittaro)
- Freeimage fixes on Windows by Christoph Gohlke
- Slightly faster _filter implementaiton

0.9.1 (August 28 2012)

- Python 3 support (you need to use 2to3)
- Haar wavelets (forward and inverse transform)
- Daubechies wavelets (forward and inverse transform)
- Corner case fix in Otsu thresholding
- Add soft_threshold function
- Have polygon.convexhull return an ndarray (instead of a list)
- Memory usage improvements in regmin/regmax/close_holes (first reported as issue #9 by thanasi)

0.9 (July 16 2012)

- Auto-convert integer to double on gaussian_filter (previously, integer values would result in zero-valued outputs).
- Check for integer types in (reg|loc)(max|min)
- Use name *out* instead of *output* for output arguments. This matches Numpy better
- Switched to MIT License

0.8.1 (June 6 2012)

- Fix gaussian_filter bug when order argument was used (reported by John Mark Agosta)
- Add morph.erode
- Improve regmax() & regmin(). Rename previous implementations to locmax() & locmin()
- Fix erode() on non-contiguous arrays

0.8 (May 7 2012)

- Move features to submodule
- Add morph.open function
- Add morph.regmax & morph.regmin functions
- Add morph.close function
- Fix morph.dilate crash

0.7.3 (March 14 2012)

- Fix installation of test data
- Greyscale erosion & dilation
- Use imread module (if available)
- Add output argument to erode() & dilate()
- Add 14th Haralick feature (patch by MattyG) — currently off by default
- Improved zernike interface (zernike_moments)
- Add remove_bordering to labeled
- Faster implementation of bwperim
- Add roundness shape feature

0.7.2 (February 13 2012)

There were two minor additions:

- Add as_rgb (especially useful for interactive use)
- Add Gaussian filtering (from scipy.ndimage)

And a few bugfixes:

- Fix type bug in 32 bit machines (Bug report by Lech Wiktor Piotrowski)
- Fix convolve1d
- Fix rank_filter

0.7.1 (January 6 2012)

The most important change fixed compilation on Mac OS X

Other changes:

- Add convolve1d
- Check that convolution arguments have right dimensions (instead of crashing)
- Add descriptor_only argument to surf.descriptors
- Specify all function signatures on freeimage.py

For version 0.7 (Dec 5 2011):

The big change was that the *dependency on scipy was removed*. As part of this process, the interpolate submodule was added. A few important bug fixes as well.

- Allow specification of centre in Zernike moment computation
- Fix Local Binary Patterns
- Remove dependency on scipy
- Add interpolate module (from scipy.ndimage)
- Add labeled_sum & labeled_sizes
- gvoronoi no longer depends on scipy

- mahotas is importable without scipy
- Fix bugs in 2D TAS (reported by Jenn Bakal)
- Support for 1-bit monochrome image loading with freeimage
- Fix GIL handling on errors (reported by Gareth McCaughan)
- Fix freeimage for 64-bit computers

Version .6.6 (August 8 2011)

- Fix fill_polygon bug (fix by joferkington)
- Fix Haralick feature 6 (fix by Rita Simões)
- Implement morph.get_structuring_element for ndim > 2. This implies that functions such as label() now also work in multiple dimensions
- Add median filter & rank_filter functions
- Add template_match function
- Refactor by use of mahotas.internal
- Better error message for when the compiled modules cannot be loaded
- Update contact email. All docs in numpydoc format now.

For version 0.6.5

- Add max_points & descriptor_only arguments to mahotas.surf
- Fix haralick for 3-D images (bug report by Rita Simões)
- Better error messages
- Fix hit&miss for non-boolean inputs
- Add label() function

For version 0.6.4:

- Fix bug in cwatershed() when using return_lines=1
- Fix bug in cwatershed() when using equivalent types for image and markers
- Move tests to mahotas.tests and include them in distribution
- Include ChangeLog in distribution
- Fix compilation on the Mac OS
- Fix compilation warnings on gcc

For version 0.6.3:

- Improve mahotas.stretch() function
- Fix corner case in surf (when determinant was zero)
- threshold argument in mahotas.surf
- imreadfromblob() & imsavetoblob() functions
- max_points argument for mahotas.surf.interest_points()

- Add `mahotas.labeled.borders` function

For version **0.6.2**:

Bugfix release:

- Fix memory leak in `_surf`
- More robust searching for freeimage
- More functions in `mahotas.surf()` to retrieve intermediate results
- Improve compilation on Windows (patches by Christoph Gohlke)

For version **0.6.1**:

- Release the GIL in morphological functions
- Convolution
- `just_filter` option in `edge.sobel()`
- `mahotas.labeled` functions
- SURF local features

For version **0.6**:

- Improve Local Binary patterns (faster and better interface)
- Much faster `erode()` (10x faster)
- Faster `dilate()` (2x faster)
- TAS for 3D images
- Haralick for 3D images

Full API Documentation

A package for computer vision in Python.

Main Features

features Compute global and local features (several submodules, include SURF and Haralick features)

convolve Convolution and wavelets

morph Morphological features. Most are available at the `mahotas` level, include `erode()`, `dilate()`...

watershed Seeded watershed implementation

imread/imsave read/write image

Documentation: <https://mahotas.readthedocs.io/>

Citation:

Coelho, Luis Pedro, 2013. Mahotas: Open source software for scriptable computer vision. Journal of Open Research Software, 1:e3, DOI: <http://dx.doi.org/10.5334/jors.ac>

`mahotas.as_rgb(r, g, b)`

Returns an RGB image with `r` in the red channel, `g` in the green, and `b` in the blue. The channels are contrast stretched.

If any of the channels is `None`, that channel is set to zero. The same can be achieved by passing `0` as that channels value. In fact, passing a number as a channel value will set the whole channel to that value.

Parameters `r,g,b` : array-like or int, optional

The channels can be of any type or `None`. At least one must be not `None` and all must have the same shape.

Returns `rgb` : ndarray

RGB ndarray

Examples

This shows a nice looking picture:

```
z1 = np.linspace(0, np.pi)
X, Y = np.meshgrid(z1, z1)
red = np.sin(X)
green = np.cos(4*Y)
blue = X*Y

plt.imshow(mahotas.as_rgb(red, green, blue))
```

Notice that the scaling on the blue channel is so different from the other channels (from 0..2500 compared with 0..1), but `as_rgb` stretches each channel independently.

`mahotas.bbox(img, border={0}, as_slice=False)`

Calculate the bounding box of image `img`.

Parameters `img` : ndarray

Any integer image type

Returns `min1,max1,min2,max2` : int,int,int,int

These are such that `img[min1:max1, min2:max2]` contains all non-zero pixels.

Returned when `as_slice` is false (the default)

`s` : slice

A slice representation of the bounding box. Returned when `as_slice` is true

`mahotas.border(labeled, i, j, Bc={3x3 cross}, out={np.zeros(labeled.shape, bool)}, always_return=True)`

Compute the border region between `i` and `j` regions.

A pixel is on the border if it has value `i` (or `j`) and a pixel in its neighbourhood (defined by `Bc`) has value `j` (or `i`).

Parameters `labeled` : ndarray of integer type

input labeled array

`i` : integer

`j` : integer

`Bc` : structure element, optional

`out` : ndarray of same shape as `labeled`, dtype=bool, optional

where to store the output. If `None`, a new array is allocated

always_return : bool, optional

if false, then, in the case where there is no pixel on the border, returns `None`. Otherwise (the default), it always returns an array even if it is empty.

Returns `border_img` : boolean ndarray

Pixels are True exactly where there is a border between i and j in *labeled*

`mahotas.borders(labeled, Bc={3x3 cross}, out={np.zeros(labeled.shape, bool)})`

Compute border pixels

A pixel is on a border if it has value i and a pixel in its neighbourhood (defined by Bc) has value j , with $i \neq j$.

Parameters `labeled` : ndarray of integer type

input labeled array

`Bc` : structure element, optional

`out` : ndarray of same shape as *labeled*, `dtype=bool`, optional

where to store the output. If `None`, a new array is allocated

`mode` : {'reflect', 'nearest', 'wrap', 'mirror', 'constant' [default], 'ignore'}

How to handle borders

Returns `border_img` : boolean ndarray

Pixels are True exactly where there is a border in *labeled*

`mahotas.bwperim(bw, n=4)`

Find the perimeter of objects in binary images.

A pixel is part of an object perimeter if its value is one and there is at least one zero-valued pixel in its neighborhood.

By default the neighborhood of a pixel is 4 nearest pixels, but if n is set to 8 the 8 nearest pixels will be considered.

Parameters `bw` : ndarray

A black-and-white image (any other image will be converted to black & white)

`n` : int, optional

Connectivity. Must be 4 or 8 (default: 4)

`mode` : {'reflect', 'nearest', 'wrap', 'mirror', 'constant' [default], 'ignore'}

How to handle borders

Returns `perim` : ndarray

A boolean image

See also:

[`borders`](#) function This is a more generic function

`mahotas.cdilate(f, g, Bc={3x3 cross}, n=1)`

Conditional dilation

cdilate creates the image y by dilating the image f by the structuring element Bc conditionally to the image g . This operator may be applied recursively n times.

Parameters **f** : Gray-scale (uint8 or uint16) or binary image.

g : Conditioning image. (Gray-scale or binary).

Bc : Structuring element (default: 3x3 cross)

n : Number of iterations (default: 1)

Returns **y** : Image

`mahotas.center_of_mass(img, labels=None)`

Returns the center of mass of img.

If *labels* is given, then it returns L centers of mass, one for each region identified by *labels* (including region 0).

Parameters **img** : ndarray

labels : ndarray, optional

A labeled array (i.e., an array of integers of the same shape as *img* such that each “object” is identified by areas with different values).

Returns **coords** : ndarray

The exact shape of the output depends on whether the *labels* argument was used.

If *labels* is None, then the return value is a 1-ndarray of coordinates (size = `len(img.shape)`); otherwise, the return value is a 2-ndarray of coordinates (shape = `(labels.max() + 1, len(img.shape))`).

`mahotas.ceroode(f, g, Bc={3x3 cross}, out={np.empty_as(A)})`

Conditional morphological erosion.

The type of operation depends on the `dtype` of A! If boolean, then the erosion is binary, else it is greyscale erosion. In the case of greyscale erosion, the smallest value in the domain of *Bc* is interpreted as -Inf.

Parameters **f** : ndarray

input image

g : ndarray

conditional image

Bc : ndarray, optional

Structuring element. By default, use a cross (see `get_structuring_elem` for details on the default).

Returns **conditionally_eroded** : ndarray

eroded version of *f* conditioned on *g*

See also:

`eroode` function Unconditional version of this function

`dilate`

`mahotas.close(f, Bc={3x3 cross}, out={np.empty_like(f)})`

Morphological closing.

close creates the image y by the morphological closing of the image f by the structuring element Bc . In the binary case, the closing by a structuring element Bc may be interpreted as the intersection of all the binary

images that contain the image f and have a hole equal to a translation of Bc . In the gray-scale case, there is a similar interpretation taking the functions umbra.

Parameters `f` : ndarray

Gray-scale (uint8 or uint16) or binary image.

`Bc` : ndarray, optional

Structuring element. (Default: 3x3 elementary cross).

`out` : ndarray, optional

Output array

`output` : deprecated

Do not use

Returns `y` : ndarray

See also:

[`open`](#) function

`mahotas.close_holes(ref, Bc=None)`
closed = close_holes(ref, Bc=None):

Close Holes

Parameters `ref` : ndarray

Reference image. This should be a binary image.

`Bc` : structuring element, optional

Default: 3x3 cross

Returns `closed` : ndarray

superset of `ref` (i.e. with closed holes)

`mahotas.convolve(f, weights, mode='reflect', cval=0.0, out={new array})`
Convolution of f and $weights$

Convolution is performed in *doubles* to avoid over/underflow, but the result is then cast to $f.dtype$. **This conversion may result in over/underflow when using small integer types or unsigned types (if the output is negative).** Converting to a floating point representation avoids this issue:

```
c = convolve(f.astype(float), kernel)
```

Parameters `f` : ndarray

input. Any dimension is supported

`weights` : ndarray

weight filter. If not of the same dtype as f , it is cast

`mode` : {'reflect' [default], 'nearest', 'wrap', 'mirror', 'constant', 'ignore'}

How to handle borders

`cval` : double, optional

If `mode` is constant, which constant to use (default: 0.0)

out : ndarray, optional

Output array. Must have same shape and dtype as f as well as be C-contiguous.

Returns convolved : ndarray of same dtype as f

`mahotas.convolve1d(f, weights, axis, mode='reflect', cval=0.0, out={new array})`

Convolution of f and $weights$ along axis $axis$.

Convolution is performed in *doubles* to avoid over/underflow, but the result is then cast to $f.dtype$.

Parameters f : ndarray

input. Any dimension is supported

weights : 1-D ndarray

weight filter. If not of the same dtype as f , it is cast

axis : int

Axis along which to convolve

mode : {'reflect' [default], 'nearest', 'wrap', 'mirror', 'constant', 'ignore'}

How to handle borders

cval : double, optional

If $mode$ is constant, which constant to use (default: 0.0)

out : ndarray, optional

Output array. Must have same shape and dtype as f as well as be C-contiguous.

Returns convolved : ndarray of same dtype as f

See also:

[**convolve**](#) function generic convolution

`mahotas.croptobbox(img, border=0)`

Returns a version of img cropped to the image's bounding box

Parameters img : ndarray

Integer image array

border : int, optional

whether to add a border (default no border)

Returns nimg : ndarray

A subimage of img .

Notes

Note that the border is on the bounding box, not on the final image! This means that if the image has a positive pixel on its margin, it will still be on the margin.

This ensures that the result is always a sub-image of the input.

```
mahotas.cwatershed(surface, markers, Bc=None, return_lines=False) W, WL = cwatershed(surface,
                                         markers, Bc=None, return_lines=True)
Seeded watershed in n-dimensions
```

This function computes the watershed transform on the input surface (which may actually be an n-dimensional volume).

This function requires initial seed points. A traditional way of initializing watershed is to use regional minima:

```
minima = mh.regmin(f)
markers, nr_markers = mh.label(minima)
W = cwatershed(f, minima)
```

Parameters `surface` : image

`markers` : image

initial markers (must be a labeled image, i.e., one where 0 represents the background and higher integers represent different regions)

`Bc` : ndarray, optional

structuring element (default: 3x3 cross)

`return_lines` : boolean, optional

whether to return separating lines (in addition to regions)

Returns `W` : integer ndarray (int64 ints)

Regions image (i.e., `W[i,j] == region for pixel (i,j)`)

`WL` : Lines image (*if return_lines==True*)

```
mahotas.daubechies(f, code, inline=False)
```

Daubechies wavelet transform

This function works best if the image sizes are powers of 2!

Parameters `f` : ndarray

2-D image

`code` : str

One of ‘D2’, ‘D4’, ... ‘D20’

`inline` : bool, optional

Whether to write the results to the input image. By default, a new image is returned.
Integer images are always converted to floating point and copied.

See also:

[`haar`](#) function Haar transform (equivalent to D2)

```
mahotas.dilate(A, Bc=None, out=None, output=None)
```

Morphological dilation.

The type of operation depends on the `dtype` of `A`! If boolean, then the dilation is binary, else it is greyscale dilation. In the case of greyscale dilation, the smallest value in the domain of `Bc` is interpreted as +Inf.

Parameters `A` : ndarray of bools

input array

Bc : ndarray, optional

Structuring element. By default, use a cross (see `get_structuring_elem` for details on the default).

out : ndarray, optional

output array. If used, this must be a C-array of the same `dtype` as `A`. Otherwise, a new array is allocated.

output : deprecated

Do not use

Returns dilated : ndarray

dilated version of `A`

See also:

`erode`

`mahotas.disk(radius, dim=2)`

Return a binary disk structuring element of radius `radius` and dimension `dim`

Parameters radius : int

Radius (in pixels) of returned disk

dim : int, optional

Dimension of returned array (default: 2)

Returns D : boolean ndarray

`mahotas.distance(bw, metric='euclidean2')`

Computes the distance transform of image `bw`:

```
dmap[i, j] = min_{i', j'} { (i-i')**2 + (j-j')**2 | !bw[i', j'] }
```

That is, at each point, compute the distance to the background.

If there is no background, then a very high value will be returned in all pixels (this is a sort of infinity).

Parameters bw : ndarray

If boolean, `False` will denote the background and `True` the foreground. If not boolean, this will be interpreted as `bw != 0` (this way you can use labeled images without any problems).

metric : str, optional

one of ‘euclidean2’ (default) or ‘euclidean’

Returns dmap : ndarray

distance map

References

For 2-D images, the following algorithm is used:

Felzenszwalb P, Huttenlocher D. *Distance transforms of sampled functions*. Cornell Computing and Information. 2004.

Available at: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.88.1647&rep=rep1&type=pdf>.

For n-D images (with $n > 2$), a slower hand-craft method is used.

`mahotas.dog(img, sigma1 = 2, thresh= None, just_filter = False)`
Compute edges using the Difference of Gaussian (DoG) operator.

edges is a binary image of edges.

Parameters `img` : Any 2D-ndarray

`sigma1` : the sigma value of the first Gaussian filter. The second filter
will have sigma value $1.001 * \text{sigma1}$

`multiplier` : the multiplier to get sigma2 . $\text{sigma2} = \text{sigma1} * \text{multiplier}$

`just_filter` : boolean, optional

If true, then return the result of filtering the image with the DoG filters, no zero-crossing
is detected (default is False).

Returns `edges` : ndarray

Binary image of edges, unless `just_filter`, in which case it will be an array of floating
point values.

`mahotas.erode(A, Bc={3x3 cross}, out={np.empty_as(A)})`
Morphological erosion.

The type of operation depends on the `dtype` of `A`! If boolean, then the erosion is binary, else it is greyscale
erosion. In the case of greyscale erosion, the smallest value in the domain of `Bc` is interpreted as -Inf.

Parameters `A` : ndarray

input image

`Bc` : ndarray, optional

Structuring element. By default, use a cross (see `get_structuring_elem` for de-
tails on the default).

`out` : ndarray, optional

output array. If used, this must be a C-array of the same `dtype` as `A`. Otherwise, a new
array is allocated.

Returns `erosion` : ndarray

eroded version of `A`

See also:

`dilate`

`mahotas.euler(f, n=8)`
Compute the Euler number of image `f`

The Euler number is also known as the Euler characteristic given that many other mathematical objects are also
known as Euler numbers.

Parameters `f` : ndarray

A 2-D binary image

`n` : int, optional

Connectivity, one of (4,8). default: 8

`mode` : {'reflect', 'nearest', 'wrap', 'mirror', 'constant' [default]}

How to handle borders

Returns `euler_nr` : int

Euler number

References

The following algorithm is used:

A Fast Algorithm for Computing the Euler Number of an Image and its VLSI Implementation, doi: 10.1109/ICVD.2000.812628

```
mahotas.find(f, template)
    Match template to image exactly
    coordinates = find(f, template)
```

The output is in the same format as the `np.where` function.

Parameters `f` : ndarray

input. Currently, only 2-dimensional images are supported.

`template` : ndarray

Template to match. Must be explicitly passed, no default.

Returns `match` : np.array

`coordinates` : np.array

These are the coordinates of the match. The format is similar to the output of `np.where`, but in an ndarray.

```
mahotas.fullhistogram(img)
```

Return a histogram with bins 0, 1, ..., “`img.max()`“.

After calling this function, it will be true that `hist[i] == (img == i).sum()`, for all `i`.

Parameters `img` : array-like of an unsigned type

input image.

Returns `hist` : an ndarray of type `np.uint32`

This will be of size `img.max() + 1`.

Notes

Only handles unsigned integer arrays.

```
mahotas.gaussian_filter(array, sigma, order=0, mode='reflect', cval=0.,
                        out=np.empty_like(array))
```

Multi-dimensional Gaussian filter.

Parameters `array` : ndarray

input array, any dimension is supported. If the array is an integer array, it will be converted to a double array.

`sigma` : scalar or sequence of scalars

standard deviation for Gaussian kernel. The standard deviations of the Gaussian filter are given for each axis as a sequence, or as a single number, in which case it is equal for all axes.

order : {0, 1, 2, 3} or sequence from same set, optional

The order of the filter along each axis is given as a sequence of integers, or as a single number. An order of 0 corresponds to convolution with a Gaussian kernel. An order of 1, 2, or 3 corresponds to convolution with the first, second or third derivatives of a Gaussian. Higher order derivatives are not implemented

mode : {'reflect' [default], 'nearest', 'wrap', 'mirror', 'constant', 'ignore'}

How to handle borders

cval : double, optional

If *mode* is constant, which constant to use (default: 0.0)

out : ndarray, optional

Output array. Must have same shape as *array* as well as be C-contiguous. If *array* is an integer array, this must be a double array; otherwise, it must have the same type as *array*.

Returns **filtered** : ndarray

Filtered version of *array*

Notes

The multi-dimensional filter is implemented as a sequence of one-dimensional convolution filters. The intermediate arrays are stored in the same data type as the output. Therefore, for output types with a limited precision, the results may be imprecise because intermediate results may be stored with insufficient precision.

```
mahotas.gaussian_filter1d(array, sigma, axis=-1, order=0, mode='reflect', cval=0.,  
                           out={np.empty_like(array)})
```

One-dimensional Gaussian filter.

Parameters **array** : ndarray

input array of a floating-point type

sigma : float

standard deviation for Gaussian kernel (in pixel units)

axis : int, optional

axis to operate on

order : {0, 1, 2, 3}, optional

An order of 0 corresponds to convolution with a Gaussian kernel. An order of 1, 2, or 3 corresponds to convolution with the first, second or third derivatives of a Gaussian. Higher order derivatives are not implemented

mode : {'reflect' [default], 'nearest', 'wrap', 'mirror', 'constant', 'ignore'}

How to handle borders

cval : double, optional

If *mode* is constant, which constant to use (default: 0.0)

out : ndarray, optional

Output array. Must have same shape and dtype as *array* as well as be C-contiguous.

Returns filtered : ndarray

Filtered version of *array*

`mahotas.get_structuring_elem(A, Bc)`

Retrieve appropriate structuring element

Parameters A : ndarray

array which will be operated on

Bc : None, int, or array-like

None Then Bc is taken to be 1

An integer

There are two associated semantics:

connectivity $Bc[y, x] = [[\text{is } |y - 1| + |x - 1| \leq Bc_i]]$

count $Bc.sum() == Bc_i$ This is the more traditional meaning (when one writes that “4-connected”, this is what one has in mind).

Fortunately, the value itself allows one to distinguish between the two semantics and, if used correctly, no ambiguity should ever occur.

An array This should be of the same nr. of dimensions as A and will be passed through if of the right type. Otherwise, it will be cast.

Returns Bc_out : ndarray

Structuring element. This array will be of the same type as A, C-contiguous.

`mahotas.haar(f, preserve_energy=True, inline=False)`

Haar transform

Parameters f : 2-D ndarray

Input image

preserve_energy : bool, optional

Whether to normalise the result so that energy is preserved (the default).

inline : bool, optional

Whether to write the results to the input image. By default, a new image is returned. Integer images are always converted to floating point and copied.

See also:

[**ihaar**](#) function Reverse Haar transform

`mahotas.hitmiss(input, Bc, out=np.zeros_like(input))`

Hit & Miss transform

For a given pixel position, the hit&miss is True if, when Bc is overlaid on *input*, centered at that position, the 1 values line up with 1s, while the 0s line up with 0s (2s correspond to *don't care*).

Parameters input : input ndarray

This is interpreted as a binary array.

Bc : ndarray

hit & miss template, values must be one of (0, 1, 2)

out : ndarray, optional

Used for output. Must be Boolean ndarray of same size as `input`

output : deprecated

Do not use

Returns `filtered` : ndarray

Examples

```
print(hitmiss(np.array([
    [0, 0, 0, 0, 0],
    [0, 1, 1, 1, 1],
    [0, 0, 1, 1, 1]]),
np.array([
    [0, 0, 0],
    [2, 1, 1],
    [2, 1, 1]])))

prints::

[[0 0 0 0 0]
 [0 0 1 1 0]
 [0 0 0 0 0]]
```

`mahotas.idaubechies(f, code, inline=False)`

Daubechies wavelet inverse transform

Parameters `f` : ndarray

2-D image

`code` : str

One of ‘D2’, ‘D4’, ... ‘D20’

`inline` : bool, optional

Whether to write the results to the input image. By default, a new image is returned. Integer images are always converted to floating point and copied.

See also:

`haar` function Haar transform (equivalent to D2)

`mahotas.ihaar(f, preserve_energy=True, inline=False)`

Reverse Haar transform

`ihhaar(haar(f))` is more or less equal to `f` (equal, except for possible rounding issues).

Parameters `f` : 2-D ndarray

Input image. If it is an integer image, it is converted to floating point (double).

`preserve_energy` : bool, optional

Whether to normalise the result so that energy is preserved (the default).

`inline` : bool, optional

Whether to write the results to the input image. By default, a new image is returned. Integer images are always converted to floating point and copied.

Returns `f` : ndarray

See also:

`haar` function Forward Haar transform

`mahotas.imread(filename, as_grey=False)`

Read an image into a ndarray from a file.

This function depends on PIL (or Pillow) being installed.

Parameters `filename` : str

filename

`as_grey` : boolean, optional

Whether to convert to grey scale image (default: no)

`mahotas.imresize(img, nsize, order=3)`

Resizes image

This function works in two ways: if `nsize` is a tuple or list of integers, then the result will be of this size; otherwise, this function behaves the same as `mh.interpolate.zoom`

Parameters `img` : ndarray

`nsize` : float or tuple(float) or tuple(integers)

Size of return. Meaning depends on the type float: `img'.shape[i] = nsize * img.shape[i]` tuple of float: `img'.shape[i] = nsize[i] * img.shape[i]` tuple of int: `img'.shape[i] = nsize[i]`

`order` : integer, optional

Spline order to use (default: 3)

Returns `img'` : ndarray

See also:

`zoom` Similar function

`scipy.misc.pilutil.imresize` Similar function

`mahotas.imwrite(filename, array)`

Writes `array` into file `filename`

This function depends on PIL (or Pillow) being installed.

Parameters `filename` : str

path on file system

`array` : ndarray-like

`mahotas.label(array, Bc={3x3 cross}, output={new array})`

Label the array, which is interpreted as a binary array

This is also called *connected component labeled*, where the connectivity is defined by the structuring element `Bc`.

See: http://en.wikipedia.org/wiki/Connected-component_labeling

Parameters `array` : ndarray

This will be interpreted as binary array

`Bc` : ndarray, optional

This is the structuring element to use

`out` : ndarray, optional

Output array. Must be a C-array, of type np.int32

Returns labeled : ndarray

Labeled result

nr_objects : int

Number of objects

`mahotas.labeled_sum(array, labeled, minlength=None)`

Labeled sum. sum will be an array of size labeled.max() + 1, where sum[i] is equal to np.sum(array[labeled == i]).

Parameters array : ndarray of any type

labeled : int ndarray

Label map. This is the same type as returned from `mahotas.label()`

minlength : int, optional

Minimum size of return array. If labeled has fewer than minlength regions, 0s are added to the result. (optional)

Returns sums : 1-d ndarray of array.dtype

`mahotas.laplacian_2D(array, alpha = 0.2)`

2D Laplacian filter.

Parameters array : ndarray

input 2D array. If the array is an integer array, it will be converted to a double array.

alpha : scalar or sequence of scalars

controls the shape of Laplacian operator. Must be 0-1. A larger values makes the operator emphasize the diagonal direction.

Returns filtered : ndarray

Filtered version of `array`

`mahotas.locmax(f, Bc={3x3 cross}, out={np.empty(f.shape, bool)})`

Local maxima

Parameters f : ndarray

Bc : ndarray, optional

structuring element

out : ndarray, optional

Used for output. Must be Boolean ndarray of same size as `f`

output : deprecated

Do not use

Returns filtered : ndarray

boolean image of same size as `f`.

See also:

`regmax` function Regional maxima. This is a stricter criterion than the local maxima as it takes the whole object into account and not just the neighbourhood defined by `Bc:: 0 0 0 0 0 0 2 0 0 0 0 2 0 0 0 0 3 0 0 0 0 3 0 0 0 0 0 0 0 0` The top 2 is a local maximum because it has the maximal value in its neighbourhood, but it is not a regional maximum.

`locmin` function Local minima

`mahotas.locmin(f, Bc={3x3 cross}, out={np.empty(f.shape, bool)})`

Local minima

Parameters `f` : ndarray

`Bc` : ndarray, optional

structuring element

`out` : ndarray, optional

Used for output. Must be Boolean ndarray of same size as `f`

`output` : deprecated

Do not use

Returns `filtered` : ndarray

boolean image of same size as `f`.

See also:

`locmax` function Regional maxima

`mahotas.majority_filter(img, N=3, out={np.empty(img.shape, np.bool)})`

Majority filter

`filtered[y,x]` is positive if the majority of pixels in the squared of size `N` centred on `(y,x)` are positive.

Parameters `img` : ndarray

input `img` (currently only 2-D images accepted)

`N` : int, optional

size of filter (must be odd integer), defaults to 3.

`out` : ndarray, optional

Used for output. Must be Boolean ndarray of same size as `img`

`output` : deprecated

Do not use

Returns `filtered` : ndarray

boolean image of same size as `img`.

`mahotas.mean_filter(f, Bc, mode='ignore', cval=0.0, out=None)`

Mean filter. The value at `mean[i, j]` will be the mean of the values in the neighbourhood defined by `Bc`.

Parameters `f` : ndarray

input. Any dimension is supported

`Bc` : ndarray

Defines the neighbourhood. Must be explicitly passed, no default.

`mode` : {‘reflect’, ‘nearest’, ‘wrap’, ‘mirror’, ‘constant’, ‘ignore’ [ignore]}

How to handle borders. The default is to ignore points beyond the border, so that the means computed near the border include fewer elements.

`cval` : double, optional

If `mode` is constant, which constant to use (default: 0.0)

`out` : ndarray, optional

Output array. Must be a double array with the same shape as f as well as be C-contiguous.

Returns **mean** : ndarray of type double and same shape as f

See also:

[**median_filter**](#) An alternative filtering method

`mahotas.median_filter(f, Bc='square', mode='reflect', cval=0.0, out=np.empty(f.shape, f.dtype))`

Median filter

Parameters **f** : ndarray

input. Any dimension is supported

Bc : ndarray or int, optional

Defines the neighbourhood, default is a square of side 3.

mode : {‘reflect’ [default], ‘nearest’, ‘wrap’, ‘mirror’, ‘constant’, ‘ignore’}

How to handle borders

cval : double, optional

If *mode* is constant, which constant to use (default: 0.0)

out : ndarray, optional

Output array. Must have same shape and dtype as f as well as be C-contiguous.

Returns **median** : ndarray of same type and shape as f

$\text{median}[i,j]$ is the median value of the points in f close to (i,j)

`mahotas.moments(img, p0, p1, cm=(0, 0), convert_to_float=True)`

Returns the p0-p1 moment of image *img*

The formula computed is

$\text{sum}_{ij} \{ \text{img}[i,j] (i - c0)^{*p0} (j - c1)^{*p1} \}$

where *cm* = (*c0,c1*). If *cm* is not given, then (0,0) is used.

If *image* is of an integer type, then it is internally converted to np.float64, unless *convert_to_float* is False. The reason is that, otherwise, overflow is likely except for small images. Since this conversion takes longer than the computation, you can turn it off in case you are sure that your images are small enough for overflow to be an issue. Note that no conversion is made if *img* is of any floating point type.

Parameters **img** : 2-ndarray

An 2-d ndarray

p0 : float

Power for first dimension

p1 : float

Power for second dimension

cm : (int,int), optional

center of mass (default: 0,0)

convert_to_float : boolean, optional

whether to convert to floating point (default: True)

normalize : boolean, optional

whether to normalize to size of image (default: False)

Returns moment: float

floating point number

Notes

It only works for 2-D images

`mahotas.open(f, Bc={3x3 cross}, out={np.empty_like(f)})`

Morphological opening.

`open` creates the image `y` by the morphological opening of the image `f` by the structuring element `Bc`.

In the binary case, the opening by the structuring element `Bc` may be interpreted as the union of translations of `b` included in `f`. In the gray-scale case, there is a similar interpretation taking the functions umbra.

Parameters `f`: ndarray

Gray-scale (uint8 or uint16) or binary image.

`Bc` : ndarray, optional

Structuring element (default: 3x3 elementary cross).

`out` : ndarray, optional

Output array

`output` : deprecated

Do not use

Returns `y` : ndarray

See also:

`open` function

`mahotas.otsu(img, ignore_zeros=False)`

Calculate a threshold according to the Otsu method.

Example:

```
import mahotas as mh
import mahotas.demos

im = mahotas.demos.nuclear_image()
# im is stored as RGB, let's convert to single 2D format:
im = im.max(2)

#Now, we compute Otsu:
t = mh.otsu(im)

# finally, we use the value to form a binary image:
bin = (im > t)
```

See Wikipedia for details on methods: http://en.wikipedia.org/wiki/Otsu's_method

Parameters `img` : an image as a numpy array.

This should be of an unsigned integer type.

`ignore_zeros` : Boolean

whether to ignore zero-valued pixels (default: False)

Returns `T` : integer

the threshold

`mahotas.overlay(gray, red=None, green=None, blue=None, if_gray_dtype_not_uint8='stretch')`

Create an image which is greyscale, but with possible boolean overlays.

Parameters `gray`: ndarray of type np.uint8

Should be a greyscale image of type np.uint8

`red,green,blue` : ndarray, optional

boolean arrays

`if_gray_dtype_not_uint8` : str, optional

What to do if `gray` is not of type np.uint8, must be one of ‘stretch’ (default): the function `stretch` is called. ‘error’ : in this case, an error is raised

Returns `overlaid` : ndarray

Colour image

`mahotas.rank_filter(f, Bc, rank, mode='reflect', cval=0.0, out=None)`

Rank filter. The value at `ranked[i, j]` will be the `rank`th largest in the neighbourhood defined by `Bc`.

Parameters `f`: ndarray

input. Any dimension is supported

`Bc` : ndarray

Defines the neighbourhood. Must be explicitly passed, no default.

`rank` : integer

`mode` : {‘reflect’ [default], ‘nearest’, ‘wrap’, ‘mirror’, ‘constant’, ‘ignore’}

How to handle borders

`cval` : double, optional

If `mode` is constant, which constant to use (default: 0.0)

`out` : ndarray, optional

Output array. Must have same shape and dtype as `f` as well as be C-contiguous.

Returns `ranked` : ndarray of same type and shape as `f`

`ranked[i,j]` is the `rank`th value of the points in `f` close to `(i,j)`

See also:

[`median_filter`](#) A special case of `rank_filter`

`mahotas.rcc(img, ignore_zeros=False)`

Calculate a threshold according to the Riddler-Calvard method.

Example:

```
import mahotas as mh
import mahotas.demos

im = mahotas.demos.nuclear_image()
# im is stored as RGB, let's convert to single 2D format:
im = im.max(2)
```

```
#Now, we compute a threshold:  
t = mh.rc(im)  
  
# finally, we use the value to form a binary image:  
bin = (im > t)
```

Parameters `img` : ndarray

Image of any type

ignore_zeros : boolean, optional

Whether to ignore zero valued pixels (default: False)

Returns `T` : float

threshold

`mahotas.regmax(f, Bc={3x3 cross}, out={np.empty(f.shape, bool)})`

Regional maxima. This is a stricter criterion than the local maxima as it takes the whole object into account and not just the neighbourhood defined by `Bc`:

```
0 0 0 0 0  
0 0 2 0 0  
0 0 2 0 0  
0 0 3 0 0  
0 0 3 0 0  
0 0 0 0 0
```

The top 2 is a local maximum because it has the maximal value in its neighbourhood, but it is not a regional maximum.

Parameters `f` : ndarray

`Bc` : ndarray, optional

structuring element

`out` : ndarray, optional

Used for output. Must be Boolean ndarray of same size as `f`

`output` : deprecated

Do not use

Returns `filtered` : ndarray

boolean image of same size as `f`.

See also:

`locmax` function Local maxima. The local maxima are a superset of the regional maxima

`mahotas.regmin(f, Bc={3x3 cross}, out={np.empty(f.shape, bool)})`

Regional minima. See the documentation for `regmax` for more details.

Parameters `f` : ndarray

`Bc` : ndarray, optional

structuring element

`out` : ndarray, optional

Used for output. Must be Boolean ndarray of same size as `f`

`output` : deprecated

Do not use

Returns **filtered** : ndarray

boolean image of same size as *f*.

See also:

[**locmin**](#) function Local minima

`mahotas.sobel(img, just_filter=False)`

Compute edges using Sobel's algorithm

edges is a binary image of edges computed according to Sobel's algorithm.

This implementation is tuned to match MATLAB's implementation.

Parameters **img** : Any 2D-ndarray

just_filter : boolean, optional

If true, then return the result of filtering the image with the sobel filters, but do not threshold (default is False).

Returns **edges** : ndarray

Binary image of edges, unless *just_filter*, in which case it will be an array of floating point values.

`mahotas.stretch(img, arg0=None, arg1=None, dtype=<type 'numpy.uint8'>)`

`img' = stretch(img, [dtype=np.uint8])` `img' = stretch(img, max, [dtype=np.uint8])` `img' = stretch(img, min, max, [dtype=np.uint8])`

Contrast stretch the image to the range [0, max] (first form) or [min, max] (second form). The method is simple linear stretching according to the formula:

```
p' = max * (p - img.min()) / img.ptp() + min
```

Parameters **img** : ndarray

input image. It is *not modified* by this function

min : integer, optional

minimum value for output [default: 0]

max : integer, optional

maximum value for output [default: 255]

dtype : dtype of output,optional

[default: np.uint8]

Returns `img'`: ndarray

resulting image. ndarray of same shape as *img* and type *dtype*.

Notes

If *max* > 255, then it truncates the values if *dtype* is not specified.

`mahotas.stretch_rgb(img, arg0=None, arg1=None, dtype=<type 'numpy.uint8'>)`

Variation of stretch() function that works per-channel on an RGB image

Parameters **img** : ndarray

input image. It is *not modified* by this function

min : integer, optional
minimum value for output [default: 0]
max : integer, optional
maximum value for output [default: 255]
dtype : dtype of output,optional
[default: np.uint8]

Returns img' : ndarray
resulting image. ndarray of same shape as *img* and type *dtype*.

See also:

[**stretch**](#) function

`mahotas.template_match(f, template, mode='reflect', cval=0.0, out=None, output=None)`

Match template to image

`match = template_match(f, template, mode='reflect', cval=0., out={np.empty_like(f)})`

The value at `match[i, j]` will be the difference (in squared euclidean terms), between *template* and a same sized window on *f* centered on that point.

Note that the computation is performed using the same dtype as *f*. Thus is may overflow if the template is large.

Parameters **f** : ndarray

input. Any dimension is supported

template : ndarray

Template to match. Must be explicitly passed, no default.

mode : {‘reflect’ [default], ‘nearest’, ‘wrap’, ‘mirror’, ‘constant’, ‘ignore’}

How to handle borders

cval : double, optional

If *mode* is constant, which constant to use (default: 0.0)

out : ndarray, optional

Output array. Must have same shape and dtype as *f* as well as be C-contiguous.

Returns **match** : ndarray of same type and shape as *f*

`match[i,j]` is the squared euclidean distance between `f[i-s0:i+s0, j-s1:j+s1]` and *template* (for appropriately defined *s0* and *s1*).

`mahotas.thin(binimg)`

Skeletonisation by thinning

Parameters **binimg** : ndarray

Binary input image

max_iter : int, optional

Maximum number of iterations (set to a negative number, the default, to run full skeletonization)

Returns **skel** : Skeletonised version of *binimg*

`mahotas.wavelet_center(f, border=0, dtype=float, cval=0.0)`

fc is a centered version of *f* with a shape that is composed of powers of 2.

Parameters `f` : ndarray
 input image
`border` : int, optional
 The border to use (default is no border)
`dtype` : type, optional
 Type of `fc`
`eval` : float, optional
 Which value to fill the border with (default is 0)

Returns `fc` : ndarray

See also:

[`wavelet_decenter`](#) function Reverse function

`mahotas.wavelet_decenter(w, oshape, border=0)`

Undoes the effect of `wavelet_center`

Parameters `w` : ndarray

 Wavelet array

`oshape` : tuple

 Desired shape

`border` : int, optional

 The desired border. This **must** be the same value as was used for `wavelet_center`

Returns `f` : ndarray

 This will have shape `oshape`

See also:

[`wavelet_center`](#) function Forward function

`mahotas.features.eccentricity(bwimage)`

Compute eccentricity

Parameters `bwimage` : ndarray

 Interpreted as a boolean image

Returns `r` : float

 Eccentricity measure

`mahotas.features.ellipse_axes(bwimage)`

Parameters of the ‘image ellipse’

`semimajor,semiminor = ellipse_axes(bwimage)`

Returns the parameters of the constant intensity ellipse with the same mass and second order moments as the original image.

Parameters `bwimage` : ndarray

 Interpreted as a boolean image

Returns `semimajor` : float

`semiminor` : float

References

Prokop, RJ, and Reeves, AP. 1992. CVGIP: Graphical Models and Image Processing 54(5):438-460

```
mahotas.features.haralick(f, ignore_zeros=False, preserve_haralick_bug=False,
                           compute_14th_feature=False, return_mean=False, re-
                           turn_mean_ptp=False, use_x_minus_y_variance=False, distance=1)
```

Compute Haralick texture features

Computes the Haralick texture features for the four 2-D directions or thirteen 3-D directions (depending on the dimensions of f).

`ignore_zeros` can be used to have the function ignore any zero-valued pixels (as background). If there are no nonzero neighbour pairs in all directions, an exception is raised. Note that this can happen even with some non-zero pixels, e.g.:

```
0 0 0 0
0 1 0 0
0 1 0 0
0 0 0 0
```

would trigger an error when `ignore_zeros=True` as there are no horizontal non-zero pairs!

Parameters `f` : ndarray of integer type

input image. 2-D and 3-D images are supported.

distance: int, optional (default=1)

The distance to consider while computing the cooccurrence matrix.

ignore_zeros : bool, optional

Whether to ignore zero pixels (default: False).

Returns `feats` : ndarray of np.double

A 4x13 or 4x14 feature vector (one row per direction) if f is 2D, 13x13 or 13x14 if it is 3D. The exact number of features depends on the value of `compute_14th_feature`. Also, if either `return_mean` or `return_mean_ptp` is set, then a single dimensional array is returned.

Other Parameters `preserve_haralick_bug` : bool, optional

whether to replicate Haralick's typo (default: False). You probably want to always set this to `False` unless you want to replicate someone else's wrong implementation.

compute_14th_feature : bool, optional

whether to compute & return the 14-th feature

return_mean : bool, optional

When set, the function returns the mean across all the directions (default: False).

return_mean_ptp : bool, optional

When set, the function returns the mean and ptp (point-to-point, i.e., difference between `max()` and `min()`) across all the directions (default: False).

use_x_minus_y_variance : bool, optional

Feature 10 (index 9) has two interpretations, as the variance of $|x-y|$ or as the variance of $P(|x-y|)$. In order to achieve compatibility with other software and

previous versions of mahotas, mahotas defaults to using $\text{VAR}[P(|x-y|)]$; if this argument is True, then it uses $\text{VAR}[|x-y|]$ (default: False)

Notes

Haralick's paper has a typo in one of the equations. This function implements the correct feature unless `preserve_haralick_bug` is True. The only reason why you'd want the buggy behaviour is if you want to match another implementation.

References

Cite the following reference for these features:

```
@article{Haralick1973,
    author = {Haralick, Robert M. and Dinstein, Its'hak and Shanmugam, K.},
    journal = {Ieee Transactions On Systems Man And Cybernetics},
    number = {6},
    pages = {610--621},
    publisher = {IEEE},
    title = {Textural features for image classification},
    url = {http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=4309314},
    volume = {3},
    year = {1973}
}
```

`mahotas.features.1bp(image, radius, points, ignore_zeros=False)`

Compute Linear Binary Patterns

The return value is a **histogram** of feature counts, where position `i` corresponds to the number of pixels that had code `i`. The codes are compressed so that impossible codes are not used. Therefore, this is the `i`'th feature, not just the feature with binary code ```i`''.

Parameters `image` : ndarray

 input image (2-D numpy ndarray)

`radius` : number (integer or floating point)

 radius (in pixels)

`points` : integer

 nr of points to consider

`ignore_zeros` : boolean, optional

 whether to ignore zeros (default: False)

Returns `features` : 1-D numpy ndarray

 histogram of features. See above for a caveat on the interpretation of these.

References

Gray Scale and Rotation Invariant Texture Classification with Local Binary Patterns Ojala, T.
Pietikainen, M. Maenpaa, T. Lecture Notes in Computer Science (Springer) 2000, ISSU 1842,
pages 404-420

`mahotas.features.pftas(img, T={mahotas.threshold.otsu(img)})`

Compute parameter free Threshold Adjacency Statistics

TAS were presented by Hamilton et al. in “Fast automated cell phenotype image classification” (<http://www.biomedcentral.com/1471-2105/8/110>)

The current version is an adapted version which is free of parameters. The thresholding is done by using Otsu’s algorithm (or can be pre-computed and passed in by setting `T`), the margin around the mean of pixels to be included is the standard deviation. This was first published by Coelho et al. in “Structured Literature Image Finder: Extracting Information from Text and Images in Biomedical Literature” (<http://www.springerlink.com/content/60634778710577t0/>)

Also returns a version computed on the negative of the binarisation defined by Hamilton et al.

Use `tas()` to get the original version of the features.

Parameters `img` : ndarray, 2D or 3D

 input image

`T` : integer, optional

 Threshold to use (default: compute with otsu)

Returns `values` : ndarray

 A 1-D ndarray of feature values

`mahotas.features.roundness(bw)`

Roundness

Parameters `bw` : ndarray

 Interpreted as a boolean image

Returns `r` : float

`mahotas.features.tas(img)`

Compute Threshold Adjacency Statistics

TAS were presented by Hamilton et al. in “Fast automated cell phenotype image classification” (<http://www.biomedcentral.com/1471-2105/8/110>)

Also returns a version computed on the negative of the binarisation defined by Hamilton et al.

See also `pftas()` for a variation without any hardcoded parameters.

Parameters `img` : ndarray, 2D or 3D

 input image

Returns `values` : ndarray

 A 1-D ndarray of feature values

See also:

`pftas` Parameter free TAS

`mahotas.features.zernike(im, degree, radius, cm={center_of_mass(im)})`

`mahotas.features.zernike_moments(im, radius, degree=8, cm={center_of_mass(im)})`

Zernike moments through `degree`. These are computed on a circle of radius `radius` centered around `cm` (or the center of mass of the image, if the `cm` argument is not used).

Returns a vector of absolute Zernike moments through `degree` for the image `im`.

Parameters `im` : 2-ndarray

 input image

radius : integer

the maximum radius for the Zernike polynomials, in pixels. Note that the area outside the circle (centered on center of mass) defined by this radius is ignored.

degree : integer, optional

Maximum degree to use (default: 8)

cm : pair of floats, optional

the centre of mass to use. By default, uses the image's centre of mass.

Returns **zvalues** : 1-ndarray of floats

Zernike moments

References

Teague, MR. (1980). Image Analysis via the General Theory of Moments. *J. Opt. Soc. Am.* 70(8):920-930.

`mahotas.colors.rgb2gray(rgb_image, dtype=np.float)`

Convert an RGB image to a grayscale image

The interpretation of RGB and greyscale values is very much object dependent (as anyone who has used an overhead projector which mangled their colour figures will have experienced). This function uses a typical method for conversion and will work acceptably well for typical use cases, but if you have strict requirements, consider implementing the conversion by yourself for fine control.

Parameters **array** : ndarray of shape (a,b,3)

dtype : dtype, optional

dtype of return

Returns **grey** : ndarray of dtype

`mahotas.colors.rgb2grey(rgb_image, dtype=np.float)`

Convert an RGB image to a grayscale image

The interpretation of RGB and greyscale values is very much object dependent (as anyone who has used an overhead projector which mangled their colour figures will have experienced). This function uses a typical method for conversion and will work acceptably well for typical use cases, but if you have strict requirements, consider implementing the conversion by yourself for fine control.

Parameters **array** : ndarray of shape (a,b,3)

dtype : dtype, optional

dtype of return

Returns **grey** : ndarray of dtype

`mahotas.colors.rgb2lab(rgb, dtype={float})`

Convert sRGB to L*a*b* coordinates

<http://en.wikipedia.org/wiki/CIELAB>

Parameters **rgb** : ndarray

Must be of shape (h,w,3)

dtype : dtype, optional

What dtype to return. Default will be floats

Returns **lab** : ndarray

`mahotas.colors.rgb2sepia(rgb)`

Parameters `rgb` : ndarray

Must be of shape (h,w,3)

Returns `sepia` : ndarray

Output is of same shape as `rgb`

`mahotas.colors.rgb2xyz(rgb, dtype={float})`

Convert RGB to XYZ coordinates

The input is interpreted as sRGB. See Wikipedia for more details:

<http://en.wikipedia.org/wiki/SRGB>

Parameters `rgb` : ndarray

`dtype` : dtype, optional

What `dtype` to return

Returns `xyz` : ndarray

See also:

`xyz2rgb` function The reverse function

`mahotas.colors.xyz2lab(xyz, dtype={float})`

Convert CIE XYZ to L*a*b* coordinates

<http://en.wikipedia.org/wiki/CIELAB>

Parameters `xyz` : ndarray

`dtype` : dtype, optional

What `dtype` to return. Default will be floats

Returns `lab` : ndarray

`mahotas.colors.xyz2rgb(xyz, dtype={float})`

Convert XYZ to sRGB coordinates

The output should be interpreted as sRGB. See Wikipedia for more details:

<http://en.wikipedia.org/wiki/SRGB>

Parameters `xyz` : ndarray

`dtype` : dtype, optional

What `dtype` to return. Default will be floats

Returns `rgb` : ndarray

See also:

`rgb2xyz` function The reverse function

CHAPTER 3

Indices and tables

- genindex
- search

Bibliography

- [R1] K. Benkrid, D. Crookes. Design and FPGA Implementation of a Perimeter Estimator. The Queen's University of Belfast. <http://www.cs.qub.ac.uk/~d.crookes/webpubs/papers/perimeter.doc>

Python Module Index

m

`mahotas`, 77

`mahotas.colors`, 103

`mahotas.features`, 99

Index

A

as_rgb() (in module mahotas), 77

B

bbox() (in module mahotas), 78
border() (in module mahotas), 78
borders() (in module mahotas), 79
bwperim() (in module mahotas), 79

C

cdilate() (in module mahotas), 79
center_of_mass() (in module mahotas), 80
cerode() (in module mahotas), 80
close() (in module mahotas), 80
close_holes() (in module mahotas), 81
convolve() (in module mahotas), 81
convolve1d() (in module mahotas), 82
cryptobox() (in module mahotas), 82
cwatershed() (in module mahotas), 82

D

daubechies() (in module mahotas), 83
dilate() (in module mahotas), 83
disk() (in module mahotas), 84
distance() (in module mahotas), 84
dog() (in module mahotas), 85

E

eccentricity() (in module mahotas.features), 99
ellipse_axes() (in module mahotas.features), 99
erode() (in module mahotas), 85
euler() (in module mahotas), 85

F

find() (in module mahotas), 86
fullhistogram() (in module mahotas), 86

G

gaussian_filter() (in module mahotas), 86

gaussian_filter1d() (in module mahotas), 87
get_structuring_elem() (in module mahotas), 88

H

haar() (in module mahotas), 88
haralick() (in module mahotas.features), 100
hitmiss() (in module mahotas), 88

I

idaubechies() (in module mahotas), 89
haar() (in module mahotas), 89
imread() (in module mahotas), 90
imresize() (in module mahotas), 90
imsave() (in module mahotas), 90

L

label() (in module mahotas), 90
labeled_sum() (in module mahotas), 91
laplacian_2D() (in module mahotas), 91
lbp() (in module mahotas.features), 101
locmax() (in module mahotas), 91
locmin() (in module mahotas), 92

M

mahotas (module), 77
mahotas.colors (module), 103
mahotas.features (module), 99
majority_filter() (in module mahotas), 92
mean_filter() (in module mahotas), 92
median_filter() (in module mahotas), 93
moments() (in module mahotas), 93

O

open() (in module mahotas), 94
otsu() (in module mahotas), 94
overlay() (in module mahotas), 95

P

pftas() (in module mahotas.features), 101

R

rank_filter() (in module mahotas), 95
rc() (in module mahotas), 95
regmax() (in module mahotas), 96
regmin() (in module mahotas), 96
rgb2gray() (in module mahotas.colors), 103
rgb2grey() (in module mahotas.colors), 103
rgb2lab() (in module mahotas.colors), 103
rgb2sepia() (in module mahotas.colors), 103
rgb2xyz() (in module mahotas.colors), 104
roundness() (in module mahotas.features), 102

S

sobel() (in module mahotas), 97
stretch() (in module mahotas), 97
stretch_rgb() (in module mahotas), 97

T

tas() (in module mahotas.features), 102
template_match() (in module mahotas), 98
thin() (in module mahotas), 98

W

wavelet_center() (in module mahotas), 98
wavelet_decenter() (in module mahotas), 99

X

xyz2lab() (in module mahotas.colors), 104
xyz2rgb() (in module mahotas.colors), 104

Z

zernike() (in module mahotas.features), 102
zernike_moments() (in module mahotas.features), 102